
BTF Viewer — Application Notes

Practical workflows for using BTFViewer to analyse RTOS scheduler behaviour and diagnose real-time performance issues. Each section is self-contained — jump to the pattern that matches your problem.

Table of Contents

Workflow	When to use it
1. First look at a new trace	Orient yourself after loading a trace for the first time
2. Measure CPU utilisation and load balance	Find overloaded / idle cores on SMP systems
3. Find the worst-case execution time (WCET)	Identify the longest task slice and its context
4. Diagnose high blocking / response time	Understand why a task waits longer than expected
5. Detect priority inversion	Spot mutex-induced priority inversion and L/M/H geometry
6. Analyse core migrations on SMP	Find ping-pong tasks and lock-boundary bounces
7. Verify task deadlines and CPU budgets	Flag executions that breach real-time requirements
8. Scope analysis to a time window	Restrict every metric to a specific region of interest
9. Compare two builds or configurations	Before/after diffing with Trace Compare
10. Monitor custom signals and code regions	Heap usage, periodic counters, interval timing

Workflow	When to use it
11. Detect tick jitter and tickless modes	Evaluate scheduler timer health
12. Check mutex and semaphore correctness	Find unpaired takes/gives and core-boundary bounces
13. Export results for reports	CSV/HTML export, PNG/SVG snapshots, headless CLI

1. First Look at a New Trace

Load and orient

```
python btf_viewer.py your-trace.btf
```

Or open via **File → Open** (Ctrl+O) / drag-and-drop. Replace `your-trace.btf` with the path to your `.btf` trace file.




Figure 1: BTF Viewer — Core View with CPU Load graph and Execution Time distribution chart

BTF Viewer in **Core View** with the CPU Load graph enabled. The right panel shows **Statistics** with Execution Time Per Slice table open; a distribution chart has been opened by clicking a task row.

What to check first

Step	Action	What to look for
1	Read the status bar	Task count, segment count, total trace span
2	Press Ctrl+0	Fit the entire trace on screen
3	Switch Task View ↔ Core View	Task View: one row per task; Core View: one row per core, expandable

Step	Action	What to look for
4	Hover a segment	Tooltip shows duration, slice index, previous/next task on that core, and the gap before this slice
5	Open Statistics panel	Summary, core utilisation, and top tasks appear immediately without placing cursors
6	Check Trace Health (TICK) section	Verify tick regularity; a warning badge signals jitter or a missed-tick estimate

Tips

- **Core View** is usually better for a first look on SMP traces: you can see which cores are busy and which are idle without needing to identify individual tasks.
- Click any core label in Core View to expand/collapse that core's per-task sub-rows. Use  **Expand All** /  **Collapse All** to toggle all at once.
- The **CPU Load Graph** (toolbar **Load** button) shows per-core utilisation as a bar chart below the timeline — enable it immediately on any multi-core trace.

2. Measure CPU Utilisation and Load Balance

Goal

Determine how much CPU time each core and each task consumes, and detect load imbalance across cores on SMP systems.

Procedure

1. Open the **Statistics** panel.
2. Expand **Core Utilisation** — one bar per core, showing active (non-IDLE, non-TICK) CPU%.

-
3. Read the **Load Balance Score** badge at the top of the section.
 - Score = $100 \times (1 - \text{Gini coefficient})$: **100** = perfect balance, **0** = one core does everything.
 - The badge turns **amber** when the population standard deviation $\sigma > 30\%$, signalling significant imbalance.
 4. Expand **Top Tasks by CPU** to see which 10 tasks consume the most CPU time.

Interpreting results

Symptom	Likely cause	What to investigate next
One core > 90 %, others idle	Tasks pinned to one core, or poor affinity configuration	Core Affinity table; verify observed execution cores against declared affinity masks
All cores high but one consistently low	A bottleneck serializes work (lock, queue)	Mutex / Semaphore and Preemption Chain tables
IDLE% unexpectedly high on all cores	Work is not being generated fast enough, or tickless idle is kicking in	Trace Health (TICK) section for tickless detection
One task dominates Top Tasks list	It may be consuming its budget; check against deadline/budget thresholds	Deadlines / CPU budget section (see §7)

Example — multi-core trace

```
python btf_viewer.py your-trace.btf
```

Open the **Statistics** panel → **Core Utilisation**. Check that each core's active CPU% is within an expected range and the **Load Balance Score** badge shows $\sigma < 30\%$

(green), indicating a well-balanced workload. **Top Tasks by CPU** reveals which tasks dominate CPU consumption.

Cursor-scoped utilisation

Place two cursors (C) around a region of interest and enable **Limit to cursor range (C1-Cn)**. Core utilisation percentages are re-computed for that window only — useful for comparing a load spike against the baseline.

3. Find the Worst-Case Execution Time (WCET)

Goal

Identify the single longest on-CPU slice for each task, navigate to it on the timeline, and understand what preceded it.

Procedure

1. Open **Statistics → Execution Time Per Slice**.
2. Click the **Max** column header to sort descending.
3. The **Max** cell for each task row is a link (dotted underline) — click it to:
 - Scroll and zoom the timeline to that slice.
 - Place a bookmark annotation with the duration.
4. Open the **distribution chart** for that task (click its row) to see the full slice-duration spread:
 - **Scatter plot**: each point is one slice; x = start time, y = duration. A single outlier far above the cluster is a WCET candidate for investigation.
 - **Histogram**: the **CDF overlay** (blue curve) shows what fraction of slices finish within a given duration — useful for checking whether the WCET is a rare tail event or part of a regular pattern.

Reading the scatter plot for WCET context

- Click any scatter point → the viewer jumps to that exact slice and adds an annotation.
- Hover the segment on the timeline to read the **gap before this slice** (scheduling latency).
- Switch to **Core View**, expand the core, to see which other tasks ran on the same core immediately before and after.

Tip — p95 vs Max

If the **Max** slice is an isolated spike far above **p95**, suspect an external interrupt or OS overhead rather than a deterministic WCET. A **p95** value that is already near your deadline budget is the more actionable number for design margins.

Example — Execution Time distribution chart

```
python btf_viewer.py your-trace.btf
```

Open **Statistics → Execution Time Per Slice**, click any task row to open its distribution chart:




Figure 2: Execution Time distribution chart — scatter plot above, adaptive histogram with CDF overlay below

The scatter plot (top) shows all slices over the trace span. The histogram (bottom) automatically selects a **log-scaled duration axis** when the duration range spans more than an order of magnitude. The blue **CDF** curve rises steeply on the left when most slices are short, then levels off as longer outliers are counted. The **p50** (green) and **p95** (orange) dashed lines align with the corresponding percentile ticks on the right axis, making it easy to read off deadline-compliance percentages.

4. Diagnose High Blocking / Response Time

Goal

Find out why a task spends a long time off-CPU between consecutive activations (scheduled-to-run but waiting).

Note: BTFViewer's **Blocking Time** is identical to Tracealyzer's **Response Time** — the off-CPU gap between the end of one task activation and the start of the next.

Procedure

1. Open **Statistics** → **Blocking Time**. Sort by **Max** or **p95** to identify tasks with long waits.
2. Click a task row to open its **blocking-time distribution chart**:
 - **Scatter**: x = when the task resumed, y = how long it waited. Clusters at certain x positions often correspond to a specific preemptor or lock contention period.
3. Click the **Max** link to jump to the worst off-CPU gap.
4. While that region is visible on the timeline, open **Statistics** → **Preemption Chain Analysis**:
 - Find the victim task in the list to see which preemptors ran on the same core during its gaps.
 - Sort by **Count** or **Total** for the victim to identify the dominant preemptor.
5. If a single preemptor dominates, investigate whether it holds a mutex during that window — open **Mutex / Semaphore** pairing to correlate object pointers.

Metric relationships




Figure 3: Diagram showing the relationship between Execution Time, Block Time, and Inter-Arrival Time for consecutive task activations

For consecutive activations of Worker_7[19]: **Execution Time** = on-CPU slice duration; **Block Time** (Response Time) = off-CPU gap from the end of one slice to the start of the next; **Inter-Arrival Time** = gap between successive slice starts = Execution Time + Block Time.

Example — Blocking Time distribution chart

```
python btf_viewer.py your-trace.btf
```

Open **Statistics** → **Blocking Time**, click any task row:




Figure 4: Blocking Time distribution chart — scatter and histogram showing off-CPU gaps

*The scatter plot reveals clusters of high blocking at specific points in time — these align with periods when a higher-priority preemptor held the core for a long stretch. Cross-reference with the **Preemption Chain** table (§4 step 4) to identify the offending task.*

Example — Preemption Chain distribution chart




Figure 5: Preemption Chain distribution chart — victim preempted by a higher-priority task

Each point represents one preemption overlap event; the y-axis shows how long the preemptor held the core during the victim's off-CPU gap. Click any scatter point to jump to the preemptor's segment on the timeline and add an annotation at that moment.

Decision table

Blocking pattern	Root cause	Fix
Blocking time $\approx$ tick period	Task is runnable but waiting for the next scheduler tick (tickless or low priority)	Raise priority, or verify the tickless idle mode configuration
Large sporadic spikes correlated with one preemptor	Priority inversion or a high-priority CPU hog	Investigate Priority Inheritance and Preemption Chain tables
Blocking time $\approx$ mutex hold time	Lock contention; victim waits while holder runs a critical section	Reduce critical-section length; check Core Bounce column in Mu-tex/Semaphore
Blocking constant across time	Periodic blocking in design — expected for periodic tasks	Use inter-arrival time to verify period; compare to deadline

5. Detect Priority Inversion

Goal

Identify episodes where a low-priority task holds a resource that blocks a high-priority waiter while medium-priority tasks run unimpeded — the classic L/M/H priority inversion geometry.

Prerequisites

Enable mutex support and task-priority tracing in your RTOS configuration. The trace library must emit `priority_inherit` and `priority_disinherit` STI events when the kernel raises or restores a task's effective priority during mutex ownership transfers.

Procedure

1. Open **Statistics → Priority Inheritance** (shown only when the trace contains priority STI events).
2. Look at the **Pattern** column:
 - **Mutex inherit**: kernel raised the holder's priority — expected, correct behaviour.
 - **L/M/H pattern** or **Mutex inherit + L/M/H**: inversion geometry detected — a medium-priority task ran while the high-priority waiter was blocked.
 - **Boost only**: manual priority-set call without kernel inheritance.
3. Click a task row to open its **boost distribution chart**:
 - Orange scatter points = boost only; red = L/M/H or mutex inherit.
 - Click a point to zoom to that episode on the timeline, highlight the task, and add an annotation.
4. On the timeline, look for the **red bottom stripe** on the task row — it marks the priority-boost window from `priority_inherit` to `priority_disinherit`.
5. Expand the core in **Core View** around the episode to see L (holder, boosted), M (medium preemptor), and H (blocked waiter) running simultaneously.

Identifying the L/M/H actors

Role	How to find it
L (Low holder)	Task with a Mutex inherit pattern and the lowest base priority in the episode
M (Medium)	Tasks whose base priority falls strictly between L's base and L's boosted peak — listed in the scatter-point tooltip
H (High waiter)	Task blocked on the same mutex object pointer — cross-check with Mutex / Semaphore pairing

Example — Priority Inversion in a multi-core trace

Open **Statistics → Priority Inheritance** on any trace that contains priority STI events. The screenshots below illustrate a classic L/M/H scenario: the low-priority mutex holder (L) is boosted to the high-priority waiter's (H) priority level so the medium-priority task (M) cannot preempt it while it holds the mutex.

Timeline — Core View expanded, zoomed to the boost window:




Figure 7: Priority boost distribution chart — mutex-inherit boost episodes plotted by start time and duration

Red points = mutex-inherit episodes; orange points = manual priority changes without kernel inheritance. The **Pattern** column classifies each episode: **Mutex inherit**, **L/M/H pattern** (a medium-priority task ran while H waited), or **Boost only** (explicit priority change).

6. Analyse Core Migrations on SMP

Goal

Identify tasks that move between cores frequently, detect lock-induced core bounces, and use the migration heatmap to find when migrations cluster in time.

Procedure

1. Open **Statistics** → **Core Migrations**.

- **Rate** = migrations per second of active time (and per scheduler tick) — a high rate relative to the task's execution frequency suggests unnecessary migration.
- **Dwell** = average on-CPU slice before migrating — short dwell + high rate = ping-pong.
- **Ping-pong** count = migrations that immediately reverse direction (A→B followed by B→A) — a strong indicator of lock-bounce or affinity issues.

2. Open the **Migration Heatmap** (toolbar **Heatmap** button, visible on multi-core traces):

- **Level 1**: directed core-pair rows × time bins. Dark cells = many migrations in that period. Migrations clustered in time indicate a periodic work-load pattern or a contention event.
- **Click a hot cell** → drill into **Level 2**: per-task sub-bins within that pair/time window.
- **Click a task cell** → zoom the timeline, place cursors, switch to Task View, and filter to that task.

3. Enable **Migrated tasks only** in the Legend filter to hide single-core tasks and focus on migrators.

4. Check **Core-Pair Migration Summary** in Statistics:

- **Lock-bounce %** = fraction of migrations where a task held a mutex while crossing core boundaries (CORE_MIGRATION_WHILE_HELD warning in Mutex/Semaphore).
- High lock-bounce % on a specific core pair points to a mutex that is frequently taken on one core and released on another.

Example — Migration Heatmap for a multi-core trace

```
python btf_viewer.py your-trace.btf
```

With the trace open, click the **Heatmap** toolbar button to open the migration heatmap.

Level 1 — core-pair overview (directed pairs × 32 time bins):




Figure 8: Migration heatmap Level 1: directed core-pair rows and time bins

Each row is a directed pair (e.g. $c_0 \rightarrow c_1$). Darker cells = more migrations in that bin. Click any dark cell to drill into Level 2.

Level 2 — per-task sub-bins (after clicking a hot cell):




Figure 9: Migration heatmap Level 2: per-task sub-bins within the selected time window

Rows are tasks that migrated on the selected pair within the chosen bin; columns are 32 sub-bins inside that bin. Click a task cell to zoom the timeline, place cursors, switch to Task View, and filter to that task.

Reducing migration overhead

Symptom	Remedy
High ping-pong between two cores	Pin the task to specific cores using the RTOS core affinity API
High lock-bounce %	Release the mutex on the same core it was taken, or increase the mutex holder's priority to finish the critical section quickly
Migrations clustered in one time bin	A burst of work on one core is pushing tasks off; check if a high-priority task dominates that period (Preemption Chain)

7. Verify Task Deadlines and CPU Budgets

Goal

Automatically flag every task execution slice that exceeds a configured per-task deadline, and every task whose CPU% exceeds a global budget threshold.

Configure thresholds

Open **Settings** (Ctrl+,) → **Analysis thresholds**:

Threshold	Unit	Example
CPU budget %	0-100 % (0 = off)	25 flags any task using more than 25 % CPU
Task deadlines	nanoseconds, one Name=ns per line	Worker[0]=500000 flags slices > 500 µs

Task names must match the **display name** shown in the label column (e.g. Worker[0], SensorTask[2]).

Reading violations

Open **Statistics → Deadlines / CPU budget**:

- **Slice over deadline** table — every slice that exceeded its per-task threshold, sorted by excess. Click any row's **Max** link to jump to that slice on the timeline.
- **CPU budget exceeded** table — tasks whose CPU% in the current scope is above the global budget.

Workflow tip — cursor-scoped deadline checking

Place cursors around a specific test phase and enable **Limit to cursor range**. The violation tables update to show only slices and CPU% within that window, letting you verify deadline compliance phase by phase without noise from the rest of the trace.

8. Scope Analysis to a Time Window

Goal

Restrict all statistics, charts, exports, and metric tables to a specific region of the trace — for example, a single test run, a load spike, or the period around a known anomaly.

Placing cursors

Action	Effect
Left-click on the timeline	Place a cursor (removes one if near an existing line)
Shift + left-click	Snap to the nearest segment boundary
C	Place a cursor at the viewport centre
Right-click → Clear all cursors	Remove all
Ctrl+R	Zoom the viewport to fit exactly between C1 and the last cursor

Example — three cursors placed on a trace

`python btf_viewer.py your-trace.btf`

Left-click three times on the timeline to place **C1**, **C2**, and **C3**:




Figure 10: Timeline with three measurement cursors C1, C2, C3 placed, showing Δ time labels between consecutive cursors

The Δ badges between consecutive cursors show elapsed time between C1-C2 and C2-C3. Enable **Limit to cursor range (C1-Cn)** in the Statistics panel to restrict all metrics to the C1-C3 window.

Enabling cursor scope

In the **Statistics** panel, check **Limit to cursor range (C1-Cn)**. All metric tables, core utilisation, top tasks, charts, and exports immediately re-compute for the C1-Cn window.

The **status bar** shows a quick min/max/avg summary of segment durations that start and end within the cursor range — useful for a fast sanity check before diving into the full statistics.

Scoping rule summary

Metric	Included in cursor range
Core & task CPU%	Overlapping active time ÷ range width
Execution time per slice	Only slices fully inside the range
Blocking time	Pairs where both slices are fully inside the range
Inter-arrival	Activations whose start time falls inside the range
Preemption chain	Overlaps counted only when the victim's gap and the preemptor are inside the range
Core migrations	Events and per-core active time with any overlap in the range

9. Compare Two Builds or Configurations

Goal

Quantitatively diff CPU%, top tasks, core migrations, blocking, preemption, and sync between two trace files — for before/after comparison or A/B configuration testing.

Procedure

1. Open both traces in separate tabs (**File** → **Open** twice, or drag-and-drop two files).
2. In the **Statistics** panel footer, click **Trace Compare...**
3. Select **Trace A** and **Trace B** from the dropdowns.
4. Optionally enable **Limit to each tab's cursor range** to compare each trace's C1-Cn window independently (requires cursors placed in each tab before opening the dialog).
5. Browse the tabs: **Summary**, **Top Tasks**, **Core Migrations**, **Blocking**, **Pre-emption**, **Sync**.
6. Export with **Export CSV** or **Export HTML** for offline reporting.

What the Δ column shows

The Δ column in each table is the raw difference $B - A$ for numeric fields. Positive Δ in CPU% means a task uses more CPU in build B; negative Δ in blocking time means the task waits less in build B — an improvement.

Headless batch comparison

```
python btf_viewer.py compare a.btf b.btf \  
  --output compare-report \  
  --format html \  
  --name-a "Before" --name-b "After"
```

Produces `compare-report.html` (and/or `.csv`) without opening the GUI — suitable for CI pipelines or automated regression reports.

10. Monitor Custom Signals and Code Regions

Goal

Embed application-level measurements (heap usage, counters, code-region timing) into the trace without modifying the scheduler hooks.

Tag events — periodic scalar values

Emit a 32-bit scalar value on any of 8 named STI channels (`tag0_event` ... `tag7_event`) from any periodic callback or task. Consult your RTOS trace library documentation for the exact API; a typical call looks like:

```
/* In any periodic hook or task */  
uint32_t value = my_get_metric(); /* heap used, stack high-water, custom  
    ↪ counter, etc. */  
trace_tag_emit(0, (int)value); /* channel 0-7, 32-bit payload */
```

In BTFViewer: - The **tag0_event** (through **tag7_event**) row appears in the label column as a waveform row. - Expand the row to see the waveform. - Open **Statistics** → **Tag Analysis** for per-channel min/avg/max/p95; click a row for a scatter + histogram chart.

Example — heap usage waveform (sample trace):

```
python btf_viewer.py your-trace.btf
```




Figure 11: `tag0_event` row expanded as a waveform showing a user metric sampled on every RTOS tick

The `tag0_event` row (expanded) shows a user-defined scalar metric sampled periodically (e.g. heap bytes in use, stack high-water mark, or any 32-bit counter). Drops indicate resource release; plateaus show steady-state consumption.

Example — Tag Analysis distribution chart:

```
python btf_viewer.py example-8cores.btf
```

Open **Statistics** → **Tag Analysis**, click the `tag0_event` row:



Figure 12: Tag value distribution chart for tag0_event in example-8cores.btf

1966 samples on tag0_event (min 8,496, avg 37,587.5, max 45,504, p95 43,904). Unlike the other metric charts, the y-axis here is the raw application value, not a duration — clusters and outliers reflect whatever firmware puts in the channel (queue depth, ADC reading, free heap, etc.).

Interval events — bracketing code regions

Bracket any code region with matching interval_start / interval_stop STI events. Consult your RTOS trace library documentation for the exact API; a typical pattern is:

```
trace_interval_start(1);    /* start of timed region; id = 1 */
do_work();
trace_interval_stop(1);     /* end of timed region; id must match */
```

In BTFViewer: - Paired spans appear as **Interval N** rows at the bottom of the timeline (horizontal task view). - Hover a bar for duration; click to add an annotation. - Open **Statistics** → **Interval Analysis** for min/avg/max/p95 per interval ID. - Click any row → distribution chart; click a scatter point → jump to that interval start and annotate.

Example — Interval Analysis distribution chart:

```
python btf_viewer.py your-trace.btf
```

Open **Statistics** → **Interval Analysis**, click any interval row:




Figure 13: Interval Analysis distribution chart — scatter of start times vs duration, histogram with CDF

Each point in the scatter represents one `interval_start` → `interval_stop` pair. The histogram shows the duration distribution of all timed code regions in scope. Click any scatter point to jump to that interval's start on the timeline.

Tip — per-task interval pairing

When the BTF note includes `tid:{task_id}` (supported by trace libraries that embed the calling task ID in interval events), start/stop events pair **per interval ID and per calling task**. This lets you use the same `id` value in multiple tasks and get separate statistics for each.

11. Detect Tick Jitter and Tickless Modes

Goal

Verify that the RTOS tick fires at the expected period and detect tickless-idle operation.

Procedure

1. Open **Statistics → Trace Health (TICK)**.
2. The section shows:
 - **Tick period** (min/avg/max from consecutive TICK STI events).
 - **Mode badge** — **TICK** (regular) or **TICKLESS** (low-power tickless mode detected from a high coefficient of variation in tick intervals).
 - **Large gaps** count and estimated missed-tick count.
3. Click the **Tick Distribution...** button (bar-chart icon, visible when ≥ 2 ticks are in scope) to open a histogram of tick intervals.

Example — Tick Distribution chart

```
python btf_viewer.py example-8cores.btf
```

Open **Statistics** → **Trace Health (TICK)**, click the **Tick Distribution...** button:



Figure 14: Tick interval distribution chart — scatter and histogram of consecutive TICK gaps in example-8cores.btf

1966 TICK events at a nominal 1.000 ms period (1000 Hz) — the coefficient of variation ($\approx 24.7\%$) exceeds the 5 % threshold, so the mode badge reads **TICKLESS**. The histogram's peaks at 1×, 2×, and 3× the nominal period show idle stretches skipping several ticks before the next one fires; the 4 largest gaps (up to 2.340 ms) are flagged as estimated missed ticks.

Interpreting results

Symptom	Interpretation
TICKLESS badge with large gaps	Tickless idle active — verify the tickless idle mode is intentionally enabled

Symptom	Interpretation
Max tick period » average	Missed ticks; the tick was delayed by a long critical section or ISR
High CoV (coefficient of variation) of tick intervals	Tick jitter — interrupt latency or scheduler interference
Missed-tick estimate > 0	Some tick events were dropped before they could be recorded — increase the ring buffer event capacity

Cursor-scoped tick health

Place cursors around a suspect region and enable **Limit to cursor range** to restrict the tick analysis to that interval — helpful when tickless periods should be confined to specific application phases.

12. Check Mutex and Semaphore Correctness

Goal

Detect unpaired takes/gives, orphaned mutex holds at trace end, cross-task gives, and lock-boundary core bounces.

Prerequisites

Enable mutex and semaphore event tracing in your RTOS trace library configuration so that take, give, create, and delete STI events are emitted for each synchronization object.

Procedure

1. Open **Statistics** → **Mutex / Semaphore pairing** (shown only when mutex/semaphore STI events are present).
2. The **main table** shows each sync object (identified by its unique object pointer) with:
 - **Hold count** — number of take/give cycles.
 - **Issue count** — number of pairing anomalies.
 - **Core bounce** — number of takes on one core with the corresponding give on a different core.
 - **Avg hold** — average time between take and give.
3. Expand **Pairing issues** sub-table to see individual anomalies:

Issue type	Meaning
Orphan give	give with no matching take — double-release or give from wrong task
Cross-task give	give by a different task than the take — valid for binary semaphores but a bug for mutexes
Unmatched take	take with no matching give — mutex held at trace end or leaked lock
CORE_MIGRATION_WHILE_HELD	Task crossed core boundary while holding the lock — source of race conditions on SMP without proper locking

4. Click any **Pairing issues** row to zoom the timeline to that event, jump to the relevant segment, and add an annotation.

Core bounce investigation

A high **Core bounce** count on a mutex indicates the lock holder frequently migrates while the mutex is taken. Cross-reference with **Core-Pair Migration Summary** to find the specific pair and then zoom to that window in the timeline to see whether a scheduling policy or affinity change would help.

13. Export Results for Reports

GUI exports

Output	How to trigger	What is included
Statistics CSV	Statistics panel → Export CSV	All metric tables for the current cursor scope
Statistics HTML	Statistics panel → Export HTML	Same tables plus Priority Inheritance boost episodes, Mutex/Semaphore hold episodes, and Interval Analysis instances
Trace Compare CSV/HTML	Trace Compare dialog → Export CSV / Export HTML	Summary, Top Tasks, Core Migrations, Blocking, Preemption, Sync diff between two traces
Timeline PNG	Toolbar Shot → Snapshot Editor → Save PNG	Current timeline viewport with optional arrow/shape annotations
Timeline SVG	File → Save as SVG (Ctrl+Shift+S)	Vector SVG of current viewport
Clipboard	File → Copy Image to Clipboard (Ctrl+Shift+C)	Raw viewport image (no editor)
Distribution chart PNG/SVG	Plot dialog → Export PNG / Export SVG	Scatter + histogram for that metric

Output	How to trigger	What is included
Migration heatmap PNG/SVG	Heatmap dialog → Export PNG / Export SVG	Current heatmap drill level

Headless CLI exports (Desktop only)

No GUI is opened; output is written directly to a file.

Trace summary (human-readable)

```
python btf_viewer.py info trace.btf
```

Trace summary (machine-readable JSON)

```
python btf_viewer.py info trace.btf --json
```

Full statistics report – HTML

```
python btf_viewer.py report trace.btf --output report.html --format html
```

Full statistics report – CSV

```
python btf_viewer.py report trace.btf --output report.csv --format csv
```

Scoped to a time window (nanoseconds)

```
python btf_viewer.py report trace.btf --output report.html --lo 100000 --hi
↪ 500000
```

Two-trace comparison – HTML

```
python btf_viewer.py compare before.btf after.btf \
  --output compare.html --format html \
  --name-a "Before" --name-b "After"
```

Core migration table as CSV

```
python btf_viewer.py migrations trace.btf -o migrations.csv
```

Automating reports in CI

Chain the headless commands in a Makefile or shell script to generate artefacts after every simulation run:

```
report.html: trace.btf
python btf_viewer.py report $< --output $@ --format html
```

```
.PHONY: report
report: report.html
```

The headless commands exit with code 0 on success and non-zero on parse error — suitable for CI pass/fail gates.

Quick-Reference: Metric to Root Cause

Observed symptom	Primary metric	Secondary metrics to check
Task misses its deadline	Execution Time Per Slice (WCET)	Preemption Chain, Blocking Time, Priority Inheritance
Task waits too long between runs	Blocking Time (Max / p95)	Preemption Chain, Mutex/Semaphore, Inter-Arrival
SMP cores unevenly loaded	Core Utilisation (Load Balance Score)	Core Migrations, Core Affinity, Preemption Chain
Lock-related slowdown	Mutex/Semaphore (Issue count, Core bounce)	Blocking Time, Preemption Chain, Core Migrations
Priority inversion suspected	Priority Inheritance (L/M/H pattern)	Mutex/Semaphore pairing, Blocking Time
Tasks migrating excessively	Core Migrations (Rate, Ping-pong)	Migration Heatmap, Core Affinity violations
Tick irregularity or missed ticks	Trace Health (TICK)	Execution Time Per Slice (long slices blocking the tick)
Memory or counter regression	Tag Analysis	Interval Analysis for the code path owning that tag
Code region too slow	Interval Analysis (WCET, p95)	Execution Time Per Slice for the owning task