

AI Assistant


BTFViewer's **AI Assistant** helps you investigate RTOS traces. It organizes measured evidence, checks possible explanations, and guides you back to the relevant region on the timeline.

Scope: The AI works with BTFViewer Findings, Statistics, timeline queries, and Trace Compare results. It does not read firmware source or ELF files. what_if results are heuristic estimates. They are not RTOS scheduler simulations or measured trace data.

Where to start

Your goal	Read
Learn the product and open the AI panel	README.md → AI Assistant
Follow a repeatable diagnosis procedure	WORKFLOWS.md
Understand a metric or Statistics page	STATISTICS.md
Configure, evaluate, or implement the AI system	This document

For a first investigation, use this sequence:



The key rule is simple: **do not move directly from a finding to a mitigation.** First define the scope, investigate the cause, and verify the evidence. Then test a change and compare the new trace.

Beginner essentials

The AI Assistant does not replace trace analysis. It helps you follow a consistent investigation process. These terms are used throughout the panel and this guide:

Term	Meaning
Analysis Finding	A deterministic rule or statistic that points to something worth checking. It is a clue, not a confirmed root cause.
Scope	The trace region used for Statistics, Findings, and AI evidence. It may be the full trace or the C1-Cn cursor range.

Term	Meaning
Filter	A task, core, or migration restriction applied inside the current Scope. A visual highlight alone is not a Filter.
Evidence	A measured value, trace event, timestamp, comparison row, or tool result that can be checked in BTFViewer.
Hypothesis	A possible explanation. It is not confirmed until the supporting evidence is checked and reasonable alternatives are tested.
Tool action	An AI request to query evidence or operate the viewer. Read-only queries run immediately; viewer-changing actions wait for Apply by default.
Baseline / Candidate	Comparable traces captured before and after a change. Trace Compare measures their differences.
What-if / Optimize	Heuristic estimates used to choose an experiment. They are not measured results.

For a first trace, open **Analysis**, select the highest-priority relevant finding, confirm it in the named Statistics section, and set C1–Cn around the incident when possible. Then open **AI Assistant** and choose **Start Investigation** or **Investigate**. Follow its evidence links, use **Verify finding**, and only then try **What-if** or **Optimize**. After changing the system, capture a new trace and use **Compare**.

Contents

User guide

1. [Overview](#)
2. [Getting started](#)
3. [Investigation workflows](#)
4. [Common use cases](#)
5. [Understanding AI results — Continue the investigation](#)
6. [Configuration, models, and privacy](#)
7. [AI tools reference](#)
8. [Viewer behavior](#)

- 9. [Troubleshooting](#)
- 10. [Opening the web app from file://](#)

Engineering reference


- 11. [CLI regression gate](#)
- 12. [Benchmark and evaluation suite — Context mode benchmarking](#)
- 13. [Investigation Case](#)
- 14. [Investigation planner](#)
- 15. [Causal and temporal engines](#)
- 16. [Implementation notes](#)
- 17. [Diagrams](#)

Context-sensitive Help can use the same fragments in both language versions: `#ai-topic-<topic-id>` for a major topic and `#ai-action-<action-id>` for a user action. Existing legacy anchors remain available for older links.

Overview

This section explains what the AI Assistant does, which evidence it uses, and its limits.

Data flow and responsibility



The AI receives structured Findings and summary metrics. It does not receive the complete raw event stream.

When more detail is needed, it requests scoped evidence through the **GUI tools**. This evidence can include:

- per-task metrics;
- timeline search results;
- correlations;
- critical paths;
- Trace Compare tables.

The AI still does not read the raw `.btf` file directly.

The AI can explain evidence, identify correlations, rank possible causes, test assumptions, and provide estimates.

Deterministic Statistics and the timeline remain the source of truth.

What the panel does

- Open the **AI Assistant** panel from the panel tabs or **Ctrl+K**. If it is hidden, enable **Settings → Panels → AI Assistant panel**.
- An empty panel shows the current **Trace**, **Scope**, and **Filters**, a question box, and actions grouped by purpose. **Start Investigation** shows the workflow line, a short blurb, and the current Finding/Scope summary, then begins a guided investigation from the available findings.
- The stepper tracks **Triage → Scope → Investigate → Verify → Experiment → Compare**. Select a completed stage to return to its output.
- **Start Investigation** stays separate from the template row. Up to five dynamic shortcuts rank recent applicable templates, then most-used, then a workflow default order (Investigate, Verify finding, Explain evidence, ...). A suggested template may be outlined from the current finding, cursors, selected task, and guide stage. Unmet Compare / SMP prerequisites show as an inline line under the chips. **More templates...** always lists every template (Analysis Findings, Explain region, Auto investigate, Compare, Report, What-if, Optimize, specialist checks, and so on). Recent/usage history is local only (Web localStorage, Desktop .rc) and is not synced across hosts; **Clear** does not erase it.
- The composer shows a collapsed **Context** one-liner (Stage · Scope · Focus · Mode · Privacy). Expand it for Trace, Scope, Filters, findings, language, endpoint, and usage. Each Web assistant reply can open **View request context**.
- The header provides **Clear**, **Language...**, and **Settings...**. **Clear** removes the conversation, usage summary, and current investigation state. **Language...** sets the reply language in the system prompt and re-states it on every user turn and tool follow-up nudge (important for smaller local models). Host wrap-ups after empty model replies also localize common Evidence titles such as Critical path.
- Select text in the conversation and right-click **Ask AI (preview...)** to send that snippet as the next question (not recorded as a template). Ask AI is enabled only when two or more words are selected. The same menu offers **Copy**, **Copy conversation**, and **Save As...**
- After a reply, continue the investigation with Evidence **[Run]** (host next steps) or conversation **[Run]** on a nextstep:{action} line. English **Next check:** prose in the reply is not a button.
- The usage bar shows **Context: Compact · 4.6k tok · 3 tools · 12s** (mode, tokens, tools, and model time). **Settings → AI → Context** chooses Compact,

Balanced (default), or Full Evidence. Confidence comes from evidence, not from the mode.

- A non-empty investigation can be restored by the viewer. Clearing the conversation also clears the saved investigation state.
- **Evidence & Validation** always shows the top-level summary (Verdict, Leading explanation, Missing evidence, Next check). Nested sections include Direct/Timeline evidence, Checks, Alternatives, Investigation details, and folds inside Investigation details. A right-side  /  control toggles Expand all / Collapse all for every nested section together (Desktop pins that control to the log viewport so wide Expand-all tables cannot push it aside). Section labels use a clear size hierarchy: panel title (12px), first-level folds such as Checks and Investigation details (12px, bordered boxes), and nested folds inside Investigation details such as Confidence evolution and Tools used (11px, indented boxes). **Save As HTML** (and Markdown/Text) keeps the same Evidence hierarchy CSS and omits tool-usage cards (Calculation / Evidence queries / Apply); Expand all still works in the saved HTML.
- Completed read-only evidence-query tool batches collapse to **Evidence queries · N completed**; pending Apply, failures, and viewer-mutating cards stay expanded. Other viewer-changing actions appear as tool cards labelled **Navigation / Scope / Filter / Annotation / Export / Calculation** and wait for **Apply** or **Skip** unless **Auto-apply GUI actions** is enabled. Tool cards sit below the written reply (amber Calculation card), not as the last line of the assistant bubble. **Undo** restores cursors, viewport, highlight, annotations, **Scope (Limit to C1-Cn)**, and **Filters**.
- What-if / Optimize results show: Simulation / estimate – not measured RTOS behavior.

Toolbar **Compare** becomes available when at least two traces are open. **Query with AI...** sends the Trace Compare tables rather than the current Findings. **Save baseline** and **Score vs baseline** use the same stored profile as `baseline_score`. **Ctrl+K** provides quick access to Analysis, AI, Compare, workspace presets, and Inspect task.

Scoping an event or region

Entry point	Scope
Analysis Findings → Investigate / Explain / Verify / Auto investigate	The selected finding and its recorded evidence
Timeline segment → Ask AI about this event	The selected task, core, and segment around <code>jump:TIME</code>
Timeline → Explain this region with AI	Available with at least two cursors; uses C1-Cn
AI panel → Explain region	Uses C1-Cn when available; otherwise uses full-trace Findings
Statistics distribution → Query with AI...	The selected task, metric, and samples displayed by the plot
Migration & Corridor Inspector → Investigate with AI	Analysis scope, selected path, ping-pong/dwell, handoff heuristic, load balance, and Inspector filters
Trace Compare → Query with AI...	The comparison tables for the two selected traces

Enable **Limit to C1-Cn** when diagnosing a phase-specific issue. The prompt then includes `Cursor region window: jump:lo ... jump:hi`, and every cited `jump:TIME` should remain inside that interval.

AI context also carries the same **Filter** and **Selection** representation shown in the status bar and Legend (Task Filter, Core Filter, Migration Filter, and current Selection). Highlight remains visual-only and is not treated as a Filter. Use **Analysis → Investigate** when you want to open the supporting Statistics section without asking the AI.

Getting started

Start here if you are using the AI Assistant for the first time.

This section explains how the AI supports common investigations. For symptom-to-metric playbooks and exact ask order, use [WORKFLOWS.md](#).

First investigation


Start with the main user actions instead of choosing individual tools. **Investigate** can select deeper evidence tools when they are needed.

Step	Action	Expected result	Check before continuing
1. Triage	Triage findings or toolbar Analysis	Ranked Critical, Warning, and Info issues	The named Statistics page shows the same issue
2. Scope	Select a finding and place or apply C1-Cn	One task, incident, or time window	Enable Limit to C1-Cn for phase-specific questions
3. Investigate	Investigate ; use Root cause when a suspect is known	Hypotheses, correlations, dependencies, and critical paths	Open the cited jump:TIME, range:L0/HI, and Statistics pages. Continue with Evidence [Run] or a conversation nextstep:{action} [Run]
4. Verify	Verify with AI... or continue the plan	Supported, rejected, or insufficient verdict	Confirm scope, task names, times, contradictions, and alternatives
5. Experiment	What-if, Optimize , or an experiment plan	Ranked estimated changes	Treat results as estimates; change the system and capture a new trace

6. Compare	Open before/after traces → Compare	Measured deltas and experiment verdict	Use equivalent workloads and comparable cursor scopes
-------------------	--	--	---

For normal use, remember only these actions: **Triage findings, Investigate, Verify with AI..., Explain region, What-if / Optimize**, and **Trace Compare**. Function names such as `correlate_events` and `rank_root_causes` are advanced and implementation references.

End-to-end flow



Do not ask for a mitigation until the timeline supports the finding.

Missing Statistics data or an incorrect scope can produce a confident-looking answer with weak evidence.

Prefer built-in templates. They already use the expected metrics and units.

Built-in actions

Choose an action by the question you need to answer. The link in the first column is stable and can be used by context-sensitive Help.

Action	Use it when	What it provides
Analysis Findings	You need a measured starting point	The current deterministic findings and their supporting Statistics sections
Triage findings	Several findings compete for attention	A priority order and the first evidence to inspect
Investigate	A finding or symptom is selected	Hypotheses, related evidence, missing checks, and a next step
Explain region	C1-Cn surrounds an incident	An explanation limited to that region, its tasks, events, and Statistics
Verify finding	A possible cause needs testing	Supporting and contradictory evidence, alternatives, and a verdict
Root cause	A likely cause or affected task is already known	A causal, correlated, or temporal chain only as far as evidence supports it
Explain finding	A finding's wording or significance is unclear	A quick, technical, or deep explanation tied to that finding
Auto investigate	You want the guided workflow to select the next checks	A staged investigation that gathers evidence and stops when more data is needed
Task profile	One task is the focus	CPU, execution tail, blocking, period, migration, synchronization, and priority context
Highest latency	Long response, blocking, dispatch, or execution episodes matter	The largest relevant tails and links to their evidence

Action	Use it when	What it provides
WCET / hot CPU	Maximum observed execution or high CPU use is suspected	Observed maxima and CPU concentration; not a proof of theoretical WCET
Migration thrash	A task moves repeatedly between cores	Migration count, rate, dwell, ping-pong, handoff heuristic (not a cache-line transfer), placement, and affinity evidence
Core balance	Core loading appears uneven	Per-core utilization, Task × Core placement, and load over time
Tick health	TICK timing or large gaps look suspicious	Regularity, tickless behavior, gap evidence, and the missed-tick estimate
Priority inversion	Blocking and priority behavior may interact	Priority boosts, L/M/H patterns, mutex evidence, and preemption checks
Deadline / budget	Tasks have timing or CPU budgets	Measured values against configured or supplied thresholds
Trace Compare	Two comparable traces are open	Measured A/B deltas and the primary regression classification
What-if	One concrete change is worth estimating	A heuristic before/after estimate that must be verified with a new trace
Optimize	A likely cause is known and several experiments are possible	Ranked mitigation experiments with evidence, expected effect, and risk

Diagnostic report	The investigation is ready to share	A structured summary of scope, findings, evidence, conclusion, alternatives, and next action
--------------------------	-------------------------------------	--

Other stable entry-point links are [#ai-action-ask_event](#), [#ai-action-query_distribution](#), [#ai-action-query_compare](#), and [#ai-action-query_corridor](#):

- **Ask AI about this event** uses one selected timeline segment.
- **Query with AI...** from a distribution plot uses the displayed samples.
- **Query with AI...** from Trace Compare uses the selected comparison tables.
- **Investigate with AI** from the Migration & Corridor Inspector uses the migrations template with structured path, ping-pong, dwell, and handoff context. It does not filter the timeline or change cursors unless you select a viewer action.

Investigation workflows

Use these workflows after you have selected the task, finding, or time window to investigate.

Investigation workflow

Step	Template or tool	Why
1	Triage findings / detect_anomalies	Rank Critical / Warning / Info; open Timeline Anomalies / Worst Events / Task Health
2	Investigate / investigate — or Findings Auto investigate...	Root-cause chain, hypotheses, alternatives, suggested tools
3	Verify finding / Findings Verify with AI...	Confirmed / Rejected / Inconclusive with jump:TIME evidence

Step	Template or tool	Why
4	<code>correlate_events + query_raw_metric</code>	Merge blocking / execution / migrations / sync / PI for one task
5	<code>find_critical_path / detect_priority_inversion</code>	Preempt/block path; L/M/H inversion suspects
6	<code>build_task_dependency_graph / analyze_temporal_causality</code>	BTF wait/preempt/migrate chain
7	<code>rank_root_causes / challenge_conclusion</code>	Rank then alternatives before <code>what_if</code>
8	<code>find_related_findings / compare_tasks</code>	Adjacent findings; side-by-side task deltas
9	<code>set_cursors / zoom_to_range / highlight_task / bookmark_finding</code>	Narrow the timeline (Apply cursors unless auto-apply is on); click <code>range:L0/HI / btfrange:</code> on critical-path evidence
10	Evidence & Validation panel	Status, direct-evidence table, checks, missing evidence; [Run] on ► Next check sends a host next-check prompt. Conversation [Run] requires <code>nextstep:{action}</code>
11	<code>generate_report / close_investigation / export_investigation</code>	Structured close-out; optional <code>export_report</code>

Root cause walks deadline/WCET → preemption → blocking → mutex → inheritance → migration for the top finding. Use it when triage already named a suspect task.

Auto investigate chains verify-style steps for one finding (investigate → correlate → critical path / graph / temporal → rank → challenge → what-if) and advances the Investigation plan checklist. Use **Verify** when you already have a finding id and want a short verdict.

Explain region and Ask event

Entry	When it appears	Scope
Timeline → Explain this region with AI	Only with ≥2 cursors ; grayed when AI is off	C1-Cn; prompt gets Cursor region window: jump:lo ... jump:hi
AI panel → Explain region	Always	Same window when ≥2 cursors; full-trace Findings if none
Timeline segment → Ask AI about this event	Segment under the pointer; grayed when AI is off	One task / core / segment around jump:TIME

Stay inside the stated window: every jump:TIME in the reply should fall between C1 and Cn (or the model should say the window has no matching evidence). Enable **Limit to C1-Cn** so Statistics / Findings / query_raw_metric match that window. Clicking a jump:TIME outside the cursors usually means the model invented or reused a full-trace time — discard it and re-ask with cursors + scoped Findings.

What-if and optimize workflow

what_if and optimize_experiment are **heuristic slice-replay** tools: they reallocate measured execution slices, scale migrations / blocking, and adjust core-util balance. They do **not** simulate an RTOS kernel or a deterministic scheduler. Every result carries a disclaimer. After a promising estimate, recommend_experiments suggests validation steps (simulation / firmware / measurement).

Goal	What to run	Typical change phrases
One concrete idea	What-if → what_if	pin CS[28] to Core_0, raise priority of Low[266], reduce mutex contention 50%
Rank several ideas	Optimize → optimize_experiment (then optional optimize for qualitative notes)	Host picks pin-to-dominant / quiet core, contention –50%, priority up, migrations –50%
Soft advice only	optimize	Finding-text mitigations without scored experiments
What to try next on the bench	recommend_experiments	Validation experiments from findings heuristics

How to read the result:

Compare baseline with simulated. Check:

- migrations;
- blocking_ns;
- load_balance_score;
- the deltas.cost ranking.

A lower cost is better in the Experiment List.

Medium confidence means the change is worth testing on real hardware. **Low confidence** usually means the change description was vague or there were too few slices.

Common use cases

These examples show how to apply the workflow to common RTOS trace problems.

Use cases

Situation	Before you ask	Template / tools	Then verify
Unknown — first look	Full-trace or cursor-scoped Findings	Triage findings → Investigate	Timeline Anomalies / Worst Events / Task Health; jump:TIME
Hottest / noisiest task	Findings name a suspect	Task profile	Period / Jitter; Task Health; Task × Core; Execution / Blocking p95/p99
Tick jitter / missed ticks	Trace Health in scope	Tick health	Tick Distribution (not Period / Jitter — that page is task inter-arrival)
Confirm one finding	Select it in Analysis Findings	Verify with AI... / Verify finding	Evidence panel; timeline
Explain a time window	≥2 cursors; Limit to C1-Cn on	Context menu or Explain region	Only jump:TIME inside C1-Cn
One segment / ISR slice	Right-click the segment	Ask AI about this event	That task's row; nearby STI
Auto walk a finding	Select finding → Auto investigate...	auto_investigate	Investigation plan + Evidence

Situation	Before you ask	Template / tools	Then verify
Migration thrash / ping-pong	Scope thrash window; Findings mention the task	Migration thrash → correlate_events → What-if pin / Optimize	Task × Core; Timeline Anomalies migration bursts; Migrations Rate/Ping; Heatmap / Chord; Core Affinity
Priority inversion / PI boost	Inheritance finding or PI episodes in scope	Priority inversion / detect_priority_inversion → find_critical_path	Priority Inheritance; Mutex hold; Waiter × Owner (heuristic handoff)
High blocking / mutex wait	Scope the stall; suspect task known	Highest latency → query_raw_metric blocking/sync → What-if contention / priority	Worst Events; Waiter × Owner; Blocking p95/p99; Mutex hold; Priority Inheritance
WCET / deadline pressure	Thresholds set in Display → Analysis	WCET / hot CPU or Deadline / budget → check_budget	Timeline Anomalies / Worst Events; Period / Jitter; Task Health; Execution Max / p95 / p99; Deadlines

Situation	Before you ask	Template / tools	Then verify
Compare two tasks	Both names known	compare_tasks	Execution / Blocking / Migrations side-by-side
Related findings	One finding selected	find_related_findings	Shared task / metric / nearby times
Load imbalance across cores	Multi-core util in Findings	Core balance → analyze_traces (multi-tab) or What-if pin to quiet core	Task × Core; Load Balance Score; Concurrent Active; Core Time Breakdown
A vs B build regression	Two tabs open	Trace Compare (toolbar Compare → Query with AI...) / compare_performance / regression_explain	Compare Summary tab; Trace Compare pages; same scope on both builds
Drift vs saved baseline	Baseline profile stored (rc / localStorage)	baseline_score	Flags z >2; re-capture if needed
Rank all open traces	≥2 loaded tabs	analyze_traces	Best tab vs Migrations / LB / missed ticks

Situation	Before you ask	Template / tools	Then verify
Write-up for a review	Cause already confirmed	Diagnostic report → generate_report → export_report / export_investigation	Saved HTML/CSV/JSON; evidence times bookmarked
CI gate vs baseline	Headless CLI	analyze with --fail-on-regression (optional --ai)	Exit code + Markdown narrative

Worked examples

Migration thrash → pin affinity

1. Place at least two cursors on the thrash window. For Statistics/Findings, enable **Limit to C1-Cn**. For the Inspector, **Follow zoom** tracks Fit vs the visible window; lock **Viewport**, or set **Analysis Scope** to **Cursor C1-Cn**.
2. Open the **Migration & Corridor Inspector**, click the heatmap path that matches the hot task (e.g. CS[22]), then run **Investigate with AI** or **Migration thrash**.
3. Ask **What-if**: *pin CS[22] to its dominant core* (or run **Optimize** for ranked candidates).
4. Read Δ migrations / Δ load_balance_score. If migrations drop but LB worsens sharply, try pin-to-quietest via optimize_experiment and compare ranks.
5. On firmware: set affinity; re-capture a .btf and toolbar **Compare** the before/after tabs.

Mutex contention → shorter critical section

1. **Highest latency** / correlate_events for the waiter; confirm hold episodes in Mutex / Priority Inheritance.
2. **What-if**: *reduce mutex contention 50% for TASK* (or another %).

3. Expect lower `blocking_ns` in the simulated payload — still an estimate. Confirm by shortening the hold in code and re-tracing.

Two builds → regression narrative

1. Open baseline and candidate as tabs; match cursor scope if needed.
2. Toolbar **Compare**, then **Query with AI...** (**Trace Compare** template), or `compare_performance` then `regression_explain`.
3. Act only on High/Medium confidence deltas that Statistics on both tabs reproduce.
4. Optional: `optimize_experiment` on the candidate's hottest task to sketch mitigations (still heuristic).

Cursor window → Explain region

1. Place **C1** / **C2** on the phase of interest (e.g. 1.060 s ... 1.120 s); enable **Limit to C1-Cn**; re-open Analysis.
2. Right-click the timeline → **Explain this region with AI** (or AI panel → **Explain region**).
3. Confirm the user turn lists Cursor region window: `jump:lo ... jump:hi`. Reject any `jump:TIME` outside that window.
4. Follow up with `correlate_events` / `query_raw_metric` on tasks named in-window; bookmark evidence times.

Priority inversion finding → Verify

1. Open Analysis Findings; select a Priority Inheritance / inversion row.
2. **Verify with AI...** (or **Auto investigate...** for a longer chain).
3. Expect `detect_priority_inversion` / `query_raw_metric` (`priority_inheritance`) / `find_critical_path`; read the Evidence score and investigation tree.
4. Click in-scope `jump:TIME` links; confirm L/M/H on the timeline and in Priority Inheritance statistics.

Simulator limits

Does	Does not
Replay measured slices / migrations / blocking gaps for the Statistics scope	Run RTOS scheduling, ISRs, or cache models
Score pin / priority / contention / migration experiments	Guarantee WCET or deadline after a firmware change
Label every result as estimate / not measured	Replace timeline verification or a new capture

Phrase changes as **pin / affinity / priority / mutex / migration** so the simulator engages; vague text falls back to a qualitative estimate (simulator: none).

Understanding AI results

AI output is an interpretation of measured evidence. Use this section to understand evidence, validation, and confidence before accepting a conclusion.

Evidence and validation

Important conclusions should include:

- evidence links such as `jump:TIME`, `range:L0/HI`, and named metrics;
- follow-up checks as a dedicated line `nextstep:{action}` (nextstep: stays English, like `jump:TIME`; the braced action uses the selected language);
- confidence: **High**, **Medium**, or **Low**;
- evidence quality: **Directly observed**, **Strong correlation**, **Possible explanation**, or **Insufficient evidence**;
- alternative explanations and what would disprove the conclusion.

Primary structured investigation replies prefer these top-level headings: **Summary**, **Evidence**, **Confidence**, and **Next check** (with optional clickable **Open Statistics** → ... links using `btfstats:section/...` when a Finding maps to a Statistics section).

Actionable Evidence links reuse the normal navigation model:

- `btfjump`: — Jump to a timestamp
- `btfrange`: / Zoom C1-Cn — place cursors and zoom the range
- `btfhighlight`: — Highlight a task
- `btfstats`: — Open the named Statistics section
- `btfnext:text/N` — Run the tagged `nextstep:{action}` sentence from the conversation
- `btfnext:run/N` — Run a host-generated next investigation prompt from Evidence (not a template)

While a request is in flight, **Stop** cancels it; Timeline stays responsive and the conversation remains visible. On failure, the prompt is restored so you can edit and **Send** again. The privacy chip shows **Local** vs **Cloud**; optional task-name redaction and sensitive-trace blocking apply before cloud send.

The Evidence & Validation panel shows:

- conclusion **Status** (Confirmed / Correlated / Suspected / Not observed / Insufficient data);
- **Finding**, a collapsible **Direct evidence** or **Timeline evidence** table, and **Interpretation**;
- collapsible **Checks** and alternative explanations; expanded **Missing evidence** and one **Next action**;
- **Investigation details** for quality band, cost, tool reasons, and trees.

The primary **Next action** is a host-generated follow-up (at most three). If the model writes no prose after tools, the next round is text-only and the assistant bubble is filled from the current case when that round is empty. The same host wrap-up is used if a later evidence tool such as `find_critical_path` is the last applied action and the model returns no final answer (or the tool-round limit is reached); that empty reply is not shown as a hard error. `find_critical_path` publishes the path table and an **Evidence graph** diagram in Evidence & Validation. After Auto investigate (or any final reply), remaining warning/error findings also become Evidence next steps even when the primary case is already verified. Extra host steps stay behind **More next steps...** When the case is complete and no remaining findings remain, the panel shows **Investigation complete** plus a stop reason instead of empty Next Steps chrome.

Continue the investigation

Use **[Run]** to keep investigating in the current Investigation Case, Context, and Scope. These prompts are not AI templates and are never recorded in template history.

Where	What [Run] sends	How the viewer finds it
Evidence & Validation ► Next check (and More next steps...)	A host-generated next-check prompt (btfnext:run/N)	Stored on the Evidence panel (at most three)
Conversation	The tagged follow-up sentence (btfnext:text/N)	A dedicated line nextstep:{action} (nextstep: stays English, like jump:TIME; the braced action uses the selected language)

1. Open cited jump:TIME / range:L0/HI links and the named Statistics page first.
2. Click Evidence **[Run]** when you want the host follow-up (coverage, remaining findings, missing evidence).
3. Click conversation **[Run]** only on a nextstep:{action} line (unbraced nextstep: ... is also accepted). The viewer shows the action plus **[Run]**, and **Open Statistics** when a Statistics section is known.
4. English **Next check:** prose in the reply is not turned into a conversation **[Run]**. Localized headings such as 下一步檢查 are not parsed either unless the model also emits nextstep:{...}.
5. Confirm the follow-up on the timeline and Statistics before treating the new reply as done.

The Evidence Quality band (under Investigation details) is a diagnostic heuristic. It is **not a probability**. Timed jump:TIME rows gathered by earlier tools (investigate, correlate_events, find_critical_path) are kept when later planner tools (rank_root_causes, challenge_conclusion, ...) publish only a verdict — otherwise Start Investigation would show Evidence Score **0%** after a strong mid-run score.

After the final reply, the host validator checks task names and timestamps. It flags unknown task names and timestamps outside the cursor window.

Prefer built-in templates. They already select the relevant metrics and Statistics pages. Use natural-language questions such as “find STI wait around TaskA” when needed; the host routes them through `search_timeline`. **Analysis Findings** can triage overall findings. **Explain finding** explains the selected Analysis Finding. Other chips: **Explain region**, **Investigate**, **Verify finding**, **Root cause**, **Trace Compare**, **Triage findings**, **Task profile**, **Diagnostic report**, **What-if**, **Optimize**, **Highest latency**, **WCET / hot CPU**, **Migration thrash**, **Core balance**, **Tick health**, **Priority inversion**, **Deadline / budget**, **Auto investigate**. Findings also offer **Save recipe...** and **Story...**

Named Statistics pages the templates cite: Timeline Anomalies, Worst Events, Period / Jitter, Unified Jitter, Recurring Patterns, Task Health, Task × Core, Waiter × Owner, Response Time, Critical Path, Preemption Matrix, Mutex Blocking, Core Utilization Over Time.

Configuration, models, and privacy

Connect an endpoint

You can use any OpenAI-compatible endpoint, including Ollama (<http://localhost:11434/v1>).

Chat requests time out after 120 seconds. **Stop** can cancel a request earlier.

Use a context window of at least **8k** when possible. This gives enough room for a full Findings card and one tool round.

If a local model has a smaller context window, use **Settings → AI → Context → Compact**. Compact reduces Findings, tool schemas, tool rows, and chat history.

The shipped Ollama default is `qwen3.5:9b`:

```
ollama pull qwen3.5:9b
```

Larger local models (qwen3.5:27b, qwen3.8:27b, and gemma4:26b) require more memory and usually run more slowly. More parameters do not guarantee a better BTFViewer investigation result. In the recorded suite, qwen3.8:27b reached the same best Overall score as qwen3.5:9b, but its mean latency was about 20 times higher. Older 7B/14B ids such as qwen2.5:7b stay optional. Avoid 3B-class models for investigations: they often skip native tools, return tool JSON as text, or fail multi-step cases.

Configuration examples are available in [examples/ai](#): [ollama.json](#), [gemini.json](#), [openai.json](#), [deepseek.json](#), [grok.json](#), and [presets.json](#).

Importing a preset fills **Settings** → **AI**, including any checkbox flags defined by the file. Review the values before saving. Each preset keeps its own base URL, model, API key, authentication mode, and TLS setting. Unknown preset names are added to the model list.

Field	Meaning
Authentication	none / API key / Sign in, per preset
Model picker	Refresh the served id list, then pick a model
Self-signed TLS	Allow self-signed TLS is intended only for a trusted private endpoint; normal certificate verification is safer

When the viewer host exposes environment variables, API keys use this precedence:

1. Key entered in **Settings** → **AI**
2. OPENAI_API_KEY
3. GEMINI_API_KEY
4. OLLAMA_API_KEY

A local Ollama endpoint normally needs no key. For a custom endpoint, enter its key in the corresponding preset. Live ai-test XML may use `<api-key env="VAR">`. See [README](#) → [API keys](#) for complete examples.

Choose a model

Capability	Small local	Local 9B+	Cloud
Basic Q&A	✓	✓	✓
Tool calling	△	✓	✓
Investigation (investigate / root-cause chain / hypotheses)	△	✓	✓
Complex reasoning (multi-step correlation, alternatives)	△	✓	✓
Large traces (big Findings card / long chat history)	△	△	✓
What-if / optimize (what_if, optimize_experiment)	✓	✓	✓

✓ reliable · △ inconsistent — works sometimes but often skips native tool calls, hallucinates numbers, or truncates on long context; always verify against the timeline before trusting a result.

The recommendations below are based on the 17-case run recorded on 2026-09-04 ([AI_BENCHMARK.md](#)). Scores and latency can change with the endpoint, hardware, model build, and dataset.

If you...	Use
Want the practical local default without an API key	qwen3.5:9b with Balanced . It passed 15/17 cases at 14.5s/case. Full evidence produced its highest Overall score, 88, but passed 14/17 cases at 16.2s/case.
Want a fast cloud response	gemini-3.5-flash-lite with Full evidence . It scored 86, passed 14/17 cases, and averaged 3.0s/case. Compact is faster (2.3s/case) at Overall 82.
Want the best result from the shipped Gemini models	gemini-3.7-flash with Full evidence . It scored 86 and passed 15/17 cases at 6.1s/case.

If you...	Use
Want the newer Gemini Flash id	gemini-3.8-flash with Balanced or Full evidence (Overall 85, 14/17 PASS, 8.2–9.1s/case). Avoid Compact for this model: Overall 68 with thin final answers.
Want a second local comparison and can accept high latency	qwen3.8:27b with Balanced . It scored 88 and passed 13/17 cases, but averaged 325.2s/case. It did not provide a consistent quality advantage over qwen3.5:9b.
Can use an optional cloud model outside the shipped suite	gpt-5.6-sol with Compact produced the highest recorded result: Overall 90, 16/17 PASS, and 10.4s/case. It is not included in the shipped benchmark configuration.
Handle confidential traces	Use local Ollama. The raw trace and extracted evidence remain on the local machine.

Small local models may skip native tool calls and emit a fenced `btftool` block instead. The viewer renders the same GUI cards either way, but investigation-heavy templates need a tool-capable model such as qwen3.5:9b to chain calls reliably.

Credential storage

The HTML viewer saves AI settings in local browser storage for convenience. Treat a saved API key as locally stored application data, not as a password vault. Do not save a long-lived key on a shared or untrusted machine. Prefer a short-lived, least-privilege key when the provider supports one; local Ollama normally requires no key.

The key is not inserted into the chat prompt. It is used only in the authentication header sent to the configured endpoint. Clear the key in **Settings → AI**, or use

Settings → Reset / clear the viewer's site data when the machine changes owner or purpose.

What leaves the machine

Sent to the configured AI endpoint	Not sent
Analysis Findings: titles, severities, task names, and heuristic text	Raw .btf or .btf.gz bytes
Metrics and tool results requested by the model	The complete unrequested event stream
Scoped timeline search and correlation results	API keys in the prompt body
Trace Compare tables when requested	—
User question and short conversation history	—

	Local Ollama	Cloud endpoint
Trace file stays local	✓	✓
Findings / metrics leave the machine	No (loopback)	Yes — to that vendor
Raw BTF uploaded	No	No
API key required	Usually no	Usually yes

Prefer local Ollama for confidential traces. Redact sensitive task names in annotations before using a cloud preset.

Context mode (token usage)

Settings → AI → Context controls how much evidence is sent with each request. Compact is the token-efficient packing mode. It reduces input tokens; Compact also caps the reply at about 300–500 tokens.

	Compact	Balanced (default)	Full evidence
Findings	Top 5 by severity	Top 12	All in scope
Tool schemas	Current stage + search / raw metric	Stage plus neighbours	Complete catalog
Tool results	10 rows; rest summarised	20 rows	40 rows
Chat history	Investigation summary + last 2 turns	Last 6 turns	Last 20 turns
Diagrams	Only if requested	When useful	When useful
What-if	Top 3 candidates	Top 5	Complete

Compact still keeps the cursor region window, real task names, jump:TIME / range:L0/HI, measurements with units, confidence / evidence quality, what-if disclaimers, and at least one alternative or falsification. If Compact omits a relevant finding, ask for a specific finding id or select a larger mode.

More context did not consistently improve the benchmark result. qwen3.5:9b had its best pass count in Balanced, while gpt-5.6-sol had its best score in Compact. qwen3.8:27b and claude-sonnet-5 also scored lower in Full evidence than in Balanced. gemini-3.8-flash recovered from Overall 68 in Compact to 85 in Balanced/Full. Treat Balanced as the general starting point, then use **--compare-context** on the intended model and workload. Select Full evidence when the investigation actually needs the additional findings, tool catalog, or history; do not assume it is always more accurate or faster.

Live ai-test defaults to Full evidence. Use **--compare-context** to measure all three modes, or **--context-mode compact** (or balanced) for a single mode. Settings → Context does not apply to the CLI scorer.

AI tools reference

The current implementation exposes 59 tools. Evidence queries, investigation-state tools, and exports run immediately. The ten viewer-changing tools wait for


Apply unless **Auto-apply GUI actions** is on: `set_cursors`, `zoom_to_range`, `highlight_task`, `set_view_mode`, `open_corridor_inspector`, `open_statistics_section`, `add_annotation`, `bookmark_finding`, `clear_marks`, and `reset_view`. Names and parameters are in [Complete tool reference](#) below.

It is easier to understand the AI tools by **purpose** than by function name.

Most users should start with the built-in templates and the Investigation plan.

The individual tool schema is mainly for advanced use, debugging, and implementation reference.

Tool mental model



Apply, Skip, and Undo

Pending viewer actions appear on an Apply card. Each line is prefixed with its class: **Navigation**, **Scope**, **Filter**, **Annotation**, **Export**, or **Calculation**.

Class	Behavior	Examples
Evidence queries / Calculation	Run immediately and return measured or derived evidence without changing the timeline view	query_raw_metric, search_timeline, investigate, correlate_events, find_critical_path, verify_claim
Investigation state and export	Run immediately; may update hypotheses, memory, experiment records, or save a file	manage_hypotheses, record_experiment_outcome, investigation_memory, close_investigation, export_report, export_investigation
Navigation	Auto-apply when the batch is navigation-only	set_cursors, zoom_to_range, highlight_task
Viewer-changing (other)	With Auto-apply GUI actions off (default), wait for Apply or Skip	set_view_mode, open_corridor_inspector, open_statistics_section, add_annotation, bookmark_finding, clear_marks, reset_view

Several viewer-changing calls may arrive in one model turn and are applied as one batch. **Undo last actions** restores zoom, view mode, highlight, inspector state, cursors, marks, **Scope (Limit to C1-Cn)**, and **task/core Filters**. Export tools use the normal file-saving behavior and do not require **Apply**.

1. Scope & navigate — “Where should I look?”

Use these tools to turn an answer into a visible timeline location.

Goal	Main tools	Result
Scope a suspected phase	set_cursors, zoom_to_range	Places C1-Cn and focuses the relevant interval
Focus a task	highlight_task	Lock-highlights the task on the timeline
Change perspective	set_view_mode	Task/Core and horizontal/vertical view
Inspect a migration corridor	open_corridor_inspector	Opens the Migration & Corridor Inspector
Show the numbers behind a claim	open_statistics_section	Opens Statistics and scrolls to a section (no scope change)
Preserve evidence	bookmark_finding, add_annotation	Adds semantic or free-text timeline marks
Clean up	clear_marks, reset_view	Clears investigation clutter or restores full-span view

Beginner usage: normally accept these as GUI cards produced by **Investigate**, **Verify**, or **Explain region** rather than calling them directly.

2. Measure & search — “What happened?”

These tools fetch deterministic evidence. They do not change the trace.

Question	Main tools	Evidence returned
Where did an event occur?	search_timeline	Matching task/STI/tag/interval/pointer/migration timestamps

Question	Main tools	Evidence returned
What are this task's measured values?	query_ raw_ metric	Scoped execution, blocking, migration, sync, PI, or findings rows
Is the distribution unusual?	analyze_ distribution	p50-p99.9, standard deviation, CV, outlier rate
Is timing periodic or jittery?	analyze_ periodicity	Expected vs observed period/jitter statistics
Does the task exceed a budget?	check_ budget	WCET/response/deadline budget comparison

These tools are the evidence layer. They should be preferred over asking the model to guess a number.

3. Investigate & correlate — “What is related?”

Use this group after Triage identifies a concrete issue.


Depth	Main tools	Purpose
Triage	detect_anomalies, cluster_findings, cluster_incidents	Rank and group Findings
Investigate	investigate, plan_investigation, suggest_scope	Build hypotheses and choose the cheapest next checks
Correlate	correlate_events, find_related_findings	Merge nearby execution/blocking/migration/sync/priority evidence
Path	find_critical_path	Walk preemption, blocking, and mutex activity around an incident
Dependency	build_task_dependency_graph	Show wait/preempt/migrate/PI relationships

Depth	Main tools	Purpose
Temporal	<code>analyze_temporal_causality</code>	Build a happens-before chain from observed evidence times
Root cause	<code>build_causal_chain</code> , <code>rank_root_causes</code>	Rank hypotheses while labeling causal/correlated/temporal relationships

Investigate / Root cause / Auto investigate orchestrate these deeper tools. Most users should not need to choose the sequence manually.

4. Verify & challenge — “Is this really the cause?”

This group prevents a plausible story from being treated as a confirmed diagnosis.



Purpose	Main tools
Check one claim	<code>verify_claim</code>
Find evidence against a hypothesis	<code>detect_contradictions</code>
Decide whether enough evidence has been gathered	<code>assess_evidence_sufficiency</code>
Force alternative explanations	<code>challenge_conclusion</code>
Track hypothesis state	<code>manage_hypotheses</code>
Inspect priority-inversion evidence	<code>detect_priority_inversion</code>

The Evidence & Validation panel complements these tools with status, a direct-evidence table, checks, missing evidence, next action, and investigation details (quality, cost, trees).

5. Compare — “What changed?”


Use these tools when two or more traces are open.

Goal	Main tools	Result
Open / obtain Trace Compare	trigger_compare	Same comparison data as toolbar Compare
Compare two builds	compare_performance	Structured metric deltas + confidence
Explain the primary regression	regression_explain	Narrative tied to A/B deltas
Localize the regression	regression_localize	Suspect task / region / mechanism
Compare two tasks	compare_tasks	Side-by-side task metrics
Compare against history	baseline_score	Drift against a stored baseline
Rank several open traces	analyze_traces	Relative scheduling behavior

A comparison is only meaningful when the two traces represent equivalent workload phases.

6. Experiment — “What should I try?”

These tools operate **after** a cause has enough evidence.



Goal	Main tools
Test one concrete idea	what_if
Rank candidate changes	optimize_experiment
Get qualitative mitigation ideas	optimize
Generate bench/firmware validation steps	recommend_experiments, generate_experiment_plan
Compare prediction with measured result	validate_experiment
Store the outcome	record_experiment_outcome

Important: what_if and optimization results are estimates, not measured scheduler behavior.

7. Report & close — “What did we learn?”

Goal	Main tools
Generate structured engineering text	generate_report
Save a diagnostic report	export_report (HTML summary first; transcript in appendix)
Save a complete Investigation Case	export_investigation
Summarize the investigation	summarize_investigation_context
Remember similar cases	investigation_memory, find_similar_investigations
Finish the Case	close_investigation

HTML export_report builds a **diagnostic report**: executive summary (status + completeness), coverage, ranked findings, observation vs interpretation, in-scope evidence table, and next action. Conversation, GUI state, out-of-window evidence, and metadata sit in an expandable **Appendix**. Export runs immediately (query/export batches need no Apply); if a turn is still in flight the HTML marks **Analysis incomplete**. Optional mode: summary (default), technical, or full (appendix conversation open).

Complete tool reference

The table below is the exhaustive schema reference. Use it when implementing, debugging, or explicitly steering tool calls.

Tool	Parameters / targets	Effect
set_cursors	timestamps (1-8 trace times)	Place cursors (enables Limit to C1-Cn when two or more)
zoom_to_range	start_time, end_time	Focus the timeline between two times

Tool	Parameters / targets	Effect
highlight_task	task_name_or_id (display name, numeric id, or merge key)	Lock-highlight a task row. Unknown names are ignored so the timeline is not dimmed. Empty string clears.
set_view_mode	mode (task / core); optional orientation	Switch Task or Core view; horizontal or vertical
open_corridor_inspector	optional core_from / core_to (Core_0, 0, c0, Core 0)	Open Migration Inspector; aliases resolve the same way
open_statistics_section	section (id such as sync, block, activation, or a page title such as Mutex Blocking)	GUI: open the Statistics panel and scroll to one section; does not change scope or cursors (Apply required)
add_annotation	time, note (≤ 240 chars)	Pin an orange timeline note at a timestamp (stays on the current right-panel tab)
query_raw_metric	task, metric (priority_ inheritance, execution, migrations, blocking, sync, findings, activation, ready_ gap, switch_reason)	Read-only: return the per-task series (or aggregate for activation / ready_ gap / switch_reason) for the current Statistics scope (up to 40 rows)

Tool	Parameters / targets	Effect
export_report	optional format (html / csv / json), optional mode (summary / technical / full)	HTML diagnostic report: executive summary, coverage, ranked findings, in-scope evidence, next action; conversation/GUI/rejected evidence in <details> appendix. Runs immediately (no Apply). Mid-flight exports still download and mark Analysis incomplete in the HTML. Strips all tool-usage cards (Calculation / Evidence queries / Apply) from the transcript. json saves a full investigation package (see export_investigation).
clear_marks	optional what (annotations / cursors / bookmarks / all / everything)	Clear AI clutter. all (default) drops annotations + cursors; everything also clears bookmarks
reset_view	(none)	Fit the timeline to the full span and clear the task highlight (marks stay)

Tool	Parameters / targets	Effect
search_timeline	query; optional mode (contains / exact / regex / sti / tags / intervals / lifecycle / pointers / migrations)	Find-panel search; returns matching timestamps (up to 40)
trigger_compare	optional tab_a / tab_b (0-based tab index or filename)	Read-only Trace Compare CSV + open the same dialog as toolbar Compare (needs two loaded tabs)
investigate	optional finding_id, depth (1-5)	Read-only: investigation graph with root-cause chain, hypotheses, ranked anomalies, suggested tools
detect_anomalies	optional limit (1-40)	Read-only: rank Analysis Findings as Critical / Warning / Info
correlate_events	task; optional around_time, window	Read-only: merge blocking / execution / migration / sync / priority / Find hits into one timeline

Tool	Parameters / targets	Effect
<code>find_critical_path</code>	task; optional timestamp, window (default 2000)	Read-only: preempt/block/mutex critical path around a timestamp; also returns a mermaid graph (graph LR), graph_nodes (id/label/kind/time), and split blocking_ steps / preemption_ steps arrays. Path steps carry start/stop; Evidence bullets use clickable range: L0/HI to zoom and place C1-C2 on the episode
<code>compare_performance</code>	optional tab_a / tab_b	Read-only: structured A vs B metric deltas + confidence (two tabs); data.regression_ type classifies the primary delta as execution / scheduling / synchronization / migration / load_ balance / unknown (legacy classification values like thrashing / load_imbalance / tick_health are preserved alongside it)

Tool	Parameters / targets	Effect
generate_report	optional report_type, finding_id	Read-only: typed engineering markdown (executive / performance / root_cause / regression / optimization / bug / ci); call export_report to save
check_budget	optional budgets, tasks	Read-only: compare per-task WCET/response/deadline metrics against budgets (host builds rows from findings when tasks omitted)
optimize	optional limit (default 5)	Read-only: evidence-backed mitigation ideas (estimate disclaimer)
regression_explain	optional tab_a / tab_b	Read-only: compare two tabs then narrate the primary regression; includes the same regression_type classification
bookmark_finding	time, kind (root_cause / evidence / correlated / reference); optional note	GUI: pin a semantic investigation annotation (Apply required)

Tool	Parameters / targets	Effect
what_if	change; optional task	Read-only: heuristic slice-replay what-if (migrations / blocking / load balance; not an RTOS kernel)
optimize_experiment	optional task, limit (1-12, default 5)	Read-only: run ranked automatic pin/priority/contention/migration experiments
analyze_traces	(none)	Read-only: rank all loaded tabs by scheduling behavior
baseline_score	optional task, baseline, snapshot	Read-only: score current per-task metrics (WCET/blocking/migrations/response) against the stored historical baseline; flags $ z > 2$
recommend_experiments	optional finding_id, task, limit (1-20, default 5)	Read-only: suggest simulation / firmware / measurement validation experiments from findings heuristics
export_investigation	optional finding_id, conclusion, tools_run, evidence_times	Download the completed investigation (finding, tools run, queries, evidence, conclusion, confidence, alternatives) as a JSON package

Tool	Parameters / targets	Effect
detect_priority_inversion	optional task, window	Read-only: scan priority-inheritance boost episodes for L/M/H inversion suspects (high/medium/low task, mutex, time, duration)
find_related_findings	optional finding_id, task, metric, window, limit (1-40, default 10)	Read-only: relate Analysis Findings by shared task, metric keyword, evidence-time proximity, or severity adjacency
compare_tasks	task_a, task_b; optional metrics	Read-only: side-by-side execution/blocking/migrations/priori inheritance delta table between two tasks
explain_finding	optional finding_id, level (quick / technical / deep)	Read-only: explain one Analysis Finding at the chosen depth (host-side; uses finding text plus hypotheses)
interpret_query	question	Read-only: turn a free-form question into an explicit investigation mode/scope before other tools run

Tool	Parameters / targets	Effect
validate_experiment	optional expected, actual (metric → signed percent)	Read-only: compare expected experiment deltas with actual A vs B / what-if results (VALIDATED / PARTIALLY VALIDATED / DISPROVED)
manage_hypotheses	hypothesis_id, status (supported / possible / rejected / need_evidence); optional reason, finding_id	Immediate: update one hypothesis in the current Investigation Case
plan_investigation	optional question, finding_id	Read-only: rank hypotheses and the cheapest tool sequence
suggest_scope	optional question	Read-only: recommend task / related tasks / time window
detect_contradictions	optional hypothesis, metrics	Read-only: verdict SUPPORTED / CONTRADICTED / INSUFFICIENT
assess_evidence_sufficiency	optional tools_run	Read-only: STOP INVESTIGATION / CONTINUE / REVISE HYPOTHESIS
cluster_findings	(none)	Read-only: group related findings into incidents

Tool	Parameters / targets	Effect
generate_fingerprint	(none)	Read-only: HIGH/MEDIUM/LOW scheduling, sync, and timing bands
find_similar_investigations	optional limit	Read-only: match fingerprint against recorded experiment outcomes
regression_localize	optional label_a, label_b	Read-only: localize A vs B inflation to a task and region
build_causal_chain	(none)	Read-only: causal / correlated / temporal edges (never silent causation)
generate_experiment_plan	optional task, limit	Read-only: ranked firmware / what-if experiments
record_experiment_outcome	optional change, predicted, actual, quality	Immediate: store an experiment outcome for later similar-case matching
analyze_temporal_causality	optional task	Read-only: happens-before chain from Findings times
build_task_dependency_graph	optional task	Read-only: BTF wait/preempt/migrate/PI graph; 2-hop neighborhood + upstream tasks
decompose_response_time	optional task	Read-only: relative delay-component shares

Tool	Parameters / targets	Effect
rank_root_causes	(none)	Read-only: rank causes from findings/hypotheses
verify_claim	claim; optional claim_type, subject, object, evidence	Read-only: SUPPORTED / PARTIAL / UNSUPPORTED
challenge_conclusion	optional conclusion	Read-only: alternatives and missing evidence
investigation_memory	optional action (recall / store), record, limit	Immediate: store or recall similar investigation cases
cluster_incidents	optional window_ns	Read-only: time-proximity incident clusters
close_investigation	optional conclusion, confidence	Immediate: close the current Investigation Case with its conclusion and confidence
analyze_distribution	optional values, metric (auto / execution / blocking / priority_inheritance / tick), task	Read-only: p50/p90/p95/p99/p99.9, stddev, CV, 3-sigma outlier rate. Statistics Query with AI... on a distribution chart harvests the open plot's samples.
analyze_periodicity	optional times, expected, source (auto / tick / sti / isr / timer / release), task, durations	Read-only: expected vs p50/p99/max, RMS and peak-to-peak jitter, kind

Tool	Parameters / targets	Effect
<code>summarize_investigation_context</code>	<code>optional conclusion,</code> <code>tools_run</code>	Read-only: compact investigation snapshot

Models without native tool calling can emit a fenced `btftool` JSON block; the viewer renders the same GUI cards. Prefer a tool-capable model for investigation-heavy workflows.

Viewer behavior

BTFViewer uses one AI workflow and one set of controls. The user-facing behavior is:

- the same six investigation stages, action templates, tools, Evidence & Validation view, and verification rules;
- endpoint, model, authentication, Context mode, privacy, and auto-apply settings under **Settings → AI**;
- report and investigation exports saved through the environment's normal download or save mechanism;
- diagrams displayed inside the conversation;
- networking and certificate rules enforced by the environment that opens the viewer.

If a local `file://` launch blocks an AI endpoint, use the development or preview server described below. This is a connection restriction, not a different AI workflow.

Troubleshooting

Symptom	Cause	Try
Web: Failed to fetch / CORS	Browser blocked a cross-origin call (file:// sends Origin: null)	Prefer <code>npm run dev / make preview</code> (both proxy Ollama), or see Opening the web app from file://
401 / 403	Missing or rejected key / origin	Settings → AI → Sign in or API key (OPENAI_API_KEY / GEMINI_API_KEY / OLLAMA_API_KEY; local Ollama needs none)
CERTIFICATE_VERIFY_FAILED / self-signed TLS	Private CA or self-signed HTTPS gateway	Trust the certificate in the operating system/browser. For a trusted private endpoint only, use Settings → AI → Allow self-signed TLS when the host supports it, or use <code>http://</code> on a protected private network.

Symptom	Cause	Try
Chat probe timed out / The read operation timed out	GET /models lists ids only; inference is slow or hung	Test connection POSTs /chat/completions (non-streaming, 120s). Warm the model (ollama run MODEL) and retry. Debug with the curl probe below; if curl hangs too, the gateway's chat upstream is stuck. Try "stream": true if non-stream never returns. Lower context length on a VRAM-tight local host.
Model not found	Typed id is not served	Refresh the Model list (or Test connection) and pick a served id from the dropdown, or ollama pull it
Gemini HTTP 400 thought_signature	Gemini 3 requires a thought blob on tool follow-ups	Retry the question — the viewer echoes Gemini thought signatures
Gemini HTTP 400 function_response.name	OpenAI-compatible follow-up with empty tool_calls[].id	The viewer fills ids and role=tool names before the next turn. Retry the case.

Symptom	Cause	Try
Empty assistant (functioncallfilter / malformedfunctioncall)	Gemini (especially Flash-Lite) emitted a tool call the API rejected	The viewer retries that turn without tools. If it still fails, switch to a fuller model (for example gemini-2.5-flash) or narrow Statistics scope.
Raw btftool JSON instead of native tool calls	Model lacks or skips function calling	The viewer renders the same cards. Select Apply or enable Auto-apply GUI actions . For reliable native calls, use a tool-capable model such as qwen3.5:9b or a supported cloud model.
Ask times out (over 120s) or stays on Waiting...	Cold start, CPU offload, or VRAM spill	Stop (composer icon), warm with ollama run MODEL, retry. Use Clear between long threads. Smaller model or shorter Statistics scope if the Findings card is huge

Symptom	Cause	Try
Later turns ignore earlier facts	Chat history exceeded the context window	Clear on the AI bar, or Analysis → Query with AI... / toolbar Compare → Query with AI... for a fresh scoped prompt
Conversation Next check: has no [Run]	The viewer only linkifies <code>nextstep:{action}</code> (or unbraced <code>nextstep: ...</code>)	Click Evidence [Run] , or ask again so the reply includes a dedicated <code>nextstep:{...}</code> line
Need raw AI request/response dumps	Debugging tool rounds / provider quirks	If the host provides Settings → AI → Log MCP messages to file , enable it only while debugging and delete <code>./ai_mcp_messages.log</code> when finished.

Test connection curl

Same body the viewer sends for **Test connection** (replace BASE, MODEL, and KEY):

```
curl -vk --max-time 180 \  
  -H "Authorization: Bearer KEY" \  
  -H "Content-Type: application/json" \  
  -d \  
    '{ "model": "MODEL", "stream": false, "messages": [{"role": "user", "content": "Reply  
    ↳ with JSON only: {\"ok\": true}"}], "max_tokens": 24}' \  
  BASE/chat/completions
```

Opening the web app from `file://`

A page opened straight from disk sends `Origin: null`, which Ollama rejects with 403 — the browser then reports only `Failed to fetch`. Serving the app over `http` avoids this entirely (`npm run dev / make preview proxy Ollama for you`), This restriction applies to a browser page opened directly from the local filesystem.

To keep using `file://`, allow every origin on the Ollama side:

```
# Server started from a terminal
OLLAMA_ORIGINS="*" ollama serve
```

```
# macOS menu-bar app (Ollama.app) – a shell variable does not reach it
launchctl setenv OLLAMA_ORIGINS "*" # then quit Ollama and reopen it
```

Verify the change took effect; expect 200 and an `Access-Control-Allow-Origin` header:

```
curl -s -D - -o /dev/null -H "Origin: null" http://localhost:11434/v1/models
↪ \
  | grep -iE "^HTTP|access-control-allow-origin"
```

If a `file://` page is still refused, list the null origin explicitly with `OLLAMA_ORIGINS="*,null"`. Note that `*` lets **any** page you visit reach your local models; undo it with `launchctl unsetenv OLLAMA_ORIGINS` when done.

CLI regression gate

The headless CI mode can compare a candidate trace to a baseline and optionally ask the configured AI for a short narrative:

Headless analyze with `--fail-on-regression`:

```
python builds/btf_viewer.py analyze candidate.btf --baseline baseline.btf
↪ --fail-on-regression
python builds/btf_viewer.py analyze candidate.btf --save-baseline
↪ /tmp/base.json
python builds/btf_viewer.py analyze candidate.btf --baseline /tmp/base.json
↪ --fail-on-regression --ai
python builds/btf_viewer.py ai-test --dataset tests/ai --fail-under 70
```

```
python builds/btf_viewer.py ai-test --config examples/ai/benchmark.xml -o
  ↪ AI_BENCHMARK.md
python builds/btf_viewer.py ai-test --config examples/ai/benchmark.xml
  ↪ --compare-context -o AI_BENCHMARK.md
python builds/btf_viewer.py ai-test --config examples/ai/benchmark.xml
  ↪ --insecure
```

Or make `-C BTFViewer ai-test (AI_DATASET, AI_FAIL_UNDER)`, make `-C BTFViewer ai-test-live (AI_CONFIG, optional AI_MODELS, writes AI\_BENCHMARK.md)`, and make `-C BTFViewer ai-test-context (same as live + --compare-context)`. Dataset, scoring rules, and context-mode flags: [Benchmark / evaluation suite](#).

See also [Export → Headless CLI](#) in the user guide.

Benchmark and evaluation suite

Offline `ai-test / runOfflineBenchmark` already ships. Live runs read **model id, base URL, TLS, and API key** from a suite XML (`--config examples/ai/benchmark.xml`) and write [AI_BENCHMARK.md](#) — a rerun **merges** into an existing file (untouched models/context-modes/cases are preserved byte-for-byte; `--replace-report` overwrites fully). Commands: [CLI regression gate](#). Default live scoring uses **Full evidence** (`--context-mode full`). **--compare-context** runs Compact, Balanced, and Full on the same cases and reports score, token totals, and latency side by side. Each live model call **retries up to 10 times, pausing 10s between attempts** on transient errors (HTTP 429/503 high demand, timeouts, empty replies); auth and not-found errors are not retried.

The capability matrix above is qualitative (small local vs 9B+ vs cloud). The suite turns those expectations into repeatable measurements: **which model is most reliable for BTF Viewer trace investigation**, not which model is largest or “smartest.”

Context mode benchmarking

Live ai-test uses the same Compact / Balanced / Full evidence packing as **Settings** → **AI** → **Context** (Findings trim, stage-filtered tool schemas, Compact reply cap). Settings in the GUI do not affect the CLI scorer.

Flag	Purpose
(default)	Full evidence — complete Findings and tool catalog
--context-mode compact	Single run in Compact (or balanced, full; comma-separated for a subset)
--compare-context	Run all three modes per model on the same cases

Each live case records **overall score**, **pass/fail**, **prompt / completion / total tokens** (summed across tool follow-ups), and **elapsed time**. With --compare-context, [AI_BENCHMARK.md](#) adds a **Context mode comparison** table per model (score vs tokens vs mean latency).

```
python builds/btf_viewer.py ai-test -c examples/ai/benchmark.xml
  ↳ --compare-context -o AI_BENCHMARK.md
make -C BTFViewer ai-test-context # AI_CONFIG, optional AI_MODELS
```

Use this when choosing a default Context setting. Compact used fewer tokens for every measured model, but it did not always reduce latency or preserve the score. Compare all three modes on the intended endpoint and workload. Details: **Context mode (token usage)**.

Scope

Keep the live set focused on:

- **Gemini cloud models**
- **Local Ollama models that are practical on a typical developer workstation**

Do **not** pick local models only because they are newest or largest. Measure models that can run alongside BTF Viewer, Ollama, and the AI context/tooling workload.

Recommended models

Local — developer workstation:

- **Qwen3.5 9B** (qwen3.5:9b) — shipped in-app default and the practical local choice. Balanced passed 15/17 cases; Full evidence produced Overall 88.
- **Qwen3.8 27B** (qwen3.8:27b) — high-latency local comparison. Balanced reached Overall 88, but averaged 325.2s/case and did not consistently outperform the 9B model.

Gemini (configurable; newer ids can be added without changing the runner):

- **Gemini 3.7 Flash** (gemini-3.7-flash) — highest pass-rate shipped Gemini reference. Full evidence reached Overall 86 with 15/17 PASS at 6.1s/case.
- **Gemini 3.8 Flash** (gemini-3.8-flash) — newer Flash id. Balanced/Full evidence reached Overall 85 with 14/17 PASS at 8.2–9.1s/case; Compact scored poorly (Overall 68).
- **Gemini 3.5 Flash-Lite** (gemini-3.5-flash-lite) — latency-focused cloud reference. Full evidence reached Overall 86 at 3.0s/case.

Shipped live suite

```
|
|— Local
|   |— Qwen3.5 9B
|   |— Qwen3.8 27B
|
|— Gemini
|   |— Gemini 3.8 Flash
|   |— Gemini 3.7 Flash
|   |— Gemini 3.5 Flash-Lite
```

The recorded results show why model size alone is not a useful selection rule. The 27B local model matched the highest local Overall score only in Balanced mode, while taking about 325 seconds per case. The 9B model reached Overall 88 in Full evidence at 16.2 seconds per case. Measure both **diagnostic quality** and **practical system performance**.

Do not hard-code the model list into the runner. Copy [examples/ai/benchmark.xml](#). For a self-signed or private-CA gateway, keep `<tls-verify>false</tls-verify>` (the suite default); public HTTPS models can override to true:

```
<ai-benchmark version="1">
  <dataset>tests/ai</dataset>
```

```
<fail-under>0</fail-under>
<output>AI_BENCHMARK.md</output>
<endpoint>
  <base-url>http://localhost:11434/v1</base-url>
  <tls-verify>>false</tls-verify>
  <timeout-s>360</timeout-s>
</endpoint>
<models>
  <model id="qwen3.5:9b"/>
  <model id="qwen3.8:27b"/>
  <model id="gemini-3.8-flash" preset="gemini">
    <base-
⇒ url>https://generativelanguage.googleapis.com/v1beta/openai</base-url>
    <tls-verify>>true</tls-verify>
    <api-key env="GEMINI_API_KEY"/>
  </model>
  <model id="gemini-3.7-flash" preset="gemini">
    <base-
⇒ url>https://generativelanguage.googleapis.com/v1beta/openai</base-url>
    <tls-verify>>true</tls-verify>
    <api-key env="GEMINI_API_KEY"/>
  </model>
  <model id="gemini-3.5-flash-lite" preset="gemini">
    <base-
⇒ url>https://generativelanguage.googleapis.com/v1beta/openai</base-url>
    <tls-verify>>true</tls-verify>
    <api-key env="GEMINI_API_KEY"/>
  </model>
</models>
</ai-benchmark>

<!-- Self-signed / private CA gateway -->
<endpoint>
  <base-url>https://llm.internal.example:8443/v1</base-url>
  <tls-verify>>false</tls-verify>
  <api-key env="GATEWAY_API_KEY"/>
</endpoint>
```

<api-key env="VAR"> reads the environment first, then any text inside the element. Omit the text (and do not commit secrets). `tls-verify false`, or `ai-test --insecure`, skips certificate checks in the headless test client. `--models id1,id2` (or make `ai-test-context AI_MODELS=id1,id2`) selects a subset of `<model>` entries. A custom suite may mark models `optional="true"` so they are skipped when

their API key is missing unless you name them in `--models / AI_MODELS`. For Ollama, list the ids you actually have pulled. Record the exact model identifier and runtime configuration.

`--only-cases id1,id2` (or make `ai-test-context AI_CASES=id1,id2`) restricts scoring to specific tests/ai dataset case ids — handy for re-testing a few cases that returned `ERROR` (transient `HTTP 429/503`) without rerunning the whole suite.

When `-o/--output` already exists, `ai-test` **merges** this run into it instead of overwriting: any model/context-mode block or offline case that was actually rerun is replaced, and every other block/case already in the file is left byte-for-byte untouched (the Comparison, Context mode comparison, and Metric breakdown tables are re-computed from the merged set). This makes narrow reruns safe, e.g. re-scoring just one model in one context mode:

```
python builds/btf_viewer.py ai-test -c examples/ai/benchmark.xml \
  --models gemini-3.7-flash --context-mode full -o AI_BENCHMARK.md
make -C BTFViewer ai-test-live AI_MODELS=gemini-3.7-flash AI_CONTEXT=full
```

Pass `--replace-report` (or `AI_REPLACE=1`) to overwrite the file fully instead of merging — useful for a fresh, full-suite run you want to replace stale results wholesale.

In-app picker (not shipped): **Settings** → **AI** → **Benchmark** with checkboxes for Gemini and local Ollama models, then **Run Benchmark**.

Dataset

`tests/ai/` holds known traces for the major diagnostic scenarios. Each case has **expected facts**, not an exact natural-language answer:

```
tests/ai/
├─ migration_thrash.btf
├─ mutex_contention.btf
├─ priority_inversion.btf
├─ deadline_miss.btf
├─ load_imbalance.btf
├─ trace_regression.btf
├─ explain_region.btf
├─ adversarial_mutex_vs_starvation.btf
├─ adversarial_exec_vs_preemption.btf
├─ adversarial_correlation_not_cause.btf
├─ adversarial_out_of_scope_time.btf
```

- |— period_jitter.btf
- |— waiter_owner_handoff.btf
- |— stats_page_next_check.btf
- |— response_vs_blocking.btf
- |— preempt_matrix_vs_chain.btf
- |— mutex_block_vs_wait_queue.btf

Adversarial cases (kind adversarial) use a decoy finding or timestamp. The obvious answer is wrong:

Case	Decoy	Actual
adversarial_mutex_vs_starvation	mutex contention	CPU starvation / preemption
adversarial_exec_vs_preemption	long execution / WCET	preemption
adversarial_correlation_not_cause	ISR caused Comm latency	correlation, no causal link
adversarial_out_of_scope_time	diagnose jump:9000	timestamp outside the cursor window
period_jitter	tick health / tickless	task inter-arrival; open Period / Jitter
waiter_owner_handoff	kernel wait-queue	heuristic mutex handoff; open Waiter × Owner
stats_page_next_check	invent detect_ timeline_ anomalies	open Timeline Anomalies / Worst Events
response_vs_blocking	Blocking Time is end-to-end response	open Response Time
preempt_matrix_vs_chain	invent detect_ preemption_ matrix	open Preemption Matrix
mutex_block_vs_wait_queue	reconstruct the kernel wait queue	open Mutex Blocking

id: migration_thrash

```
trace: migration_thrash.btf
expected:
  finding_types: [migration, load_balance]
  tasks: [CS[22]]
  evidence:
    required_metrics: [migrations]
  allowed_tools: [detect_anomalies, correlate_events, query_raw_metric]
  forbidden:
    invented_task_names: true
    out_of_scope_timestamps: true
```

That keeps scoring robust against harmless wording differences.

Evaluation metrics

Metric	What it measures
Finding identification	Did the model identify the expected problem?
Evidence accuracy	Are cited metrics/events actually present? <code>required_metrics</code> also accepts Statistics page titles (Period / Jitter, Waiter × Owner, Timeline Anomalies, ...) and common aliases (Period/Jitter, Waiter × Owner)
Timestamp validity	Are <code>jump:TIME</code> values real and in scope?
Task-name validity	Did the model use only known task names?
Tool selection	Did it call appropriate investigation tools?
Tool-chain quality	Did it gather enough evidence before concluding?
Root-cause accuracy	Does the conclusion match the expected diagnosis?
Alternative handling	Did it consider plausible alternatives?
Confidence calibration	Is confidence consistent with the available evidence?
Response completeness	Did it answer the investigation question completely?
Latency	How long did the investigation take?
Tool-call count	How many tool rounds were required?
Peak memory	How much RAM was consumed during inference?

Metric	What it measures
Time to first token (TTFT)	How quickly did the model begin responding?
Generation throughput	Sustained tokens/sec during the investigation
Investigation success rate	Percentage of cases completed correctly within the configured time/resource limit
False-causal rate	Claimed a causal link the case marks as coincidence / non-causal (0-100, higher is worse)
False-confirmation rate	Confirmed the decoy finding (trap_phrases) instead of the real cause
Unsupported-claim rate	Share of validator claims that fail task/time/scope checks
Premature-conclusion rate	High confidence or a conclusion before required tools ran

For local runs, memory and latency are first-class. A slightly more accurate model that is unusable under memory pressure should not automatically rank higher.

Level 1 — tool / evidence correctness: valid tool, parameters, task, timestamp, and scope. Isolates tool-use bugs from reasoning quality.

Level 2 — diagnostic correctness: expected vs actual diagnosis, evidence, and alternatives. A convincing explanation is not enough.

Headline score: weighted engineering score (Finding / Evidence / Tool use / Root cause / Calibration / Safety), not a probability of correctness. Keep the component scores visible. PASS is overall ≥ 70 .

Safeguards the suite must test: no invented task names, metric values, or `jump:TIME`; no timestamps outside the cursor region; no unsupported conclusions presented as confirmed; evidence must match tool results; heuristic what-if stays labeled as estimates; the model must not claim a simulation is a measured result.

Model matrix

Same suite against the shipped Gemini and local Ollama models. Recorded 2026-09-04 (17-case dataset); full case tables and Compact/Balanced numbers: [AI_BEN](#)

[CHMARK.md](#). Scores below are the **Full evidence** context mode (the live-scoring default).

Model	Category	Finding	Evidence	Root cause	Calibration	Notes
qwen3.5:9b	Local / practical	85	93	82	80	overall 88 , 16.2s/case, 14/17 PASS
qwen3.8:27b	Local / high-latency	88	94	65	80	overall 86 , 332s/case, 13/17 PASS
gemini-3.5-flash-lite	Cloud / fast	82	94	65	80	overall 86 , 3.0s/case, 14/17 PASS
gemini-3.7-flash	Cloud	91	91	59	80	overall 86 , 6.1s/case, 15/17 PASS
gemini-3.8-flash	Cloud	88	91	59	80	overall 85 , 9.1s/case, 14/17 PASS

[AI_BENCHMARK.md](#) also carries two cloud models run outside the shipped suite (--models against a private config), for reference only — they are not in examples/ai/benchmark.xml and are not part of the “Recommended models” guidance below:

Model	Category	Finding	Evidence	Root cause	Calibration	Notes
claude-sonnet-5	Cloud (optional, not shipped)	85	94	59	80	overall 82 , 15.9s/case, 12/17 PASS
gpt-5.6-sol	Cloud (optional, not shipped)	91	90	76	80	overall 88 , 9.6s/case, 15/17 PASS

Live --config runs a tool-result follow-up when the first turn is tools-only (or planning text without a Confidence line). Single-turn scores are not comparable.

The results support these practical conclusions:

- **Best practical local setup:** qwen3.5:9b with Balanced for the highest pass count, or Full evidence for the highest Overall score.
- **Fastest measured cloud setup:** gemini-3.5-flash-lite with Compact at 2.3s/case; Full evidence reached Overall 86 at 3.0s/case.
- **Best shipped Gemini result:** gemini-3.7-flash with Full evidence, Overall 86 and 15/17 PASS at 6.1s/case.
- **Gemini 3.8 Flash:** Balanced/Full evidence Overall 85; avoid Compact (Overall 68).
- **Highest result in the complete report:** optional gpt-5.6-sol with Compact, Overall 90 and 16/17 PASS.
- **More context is not automatically better:** the best mode depends on the model. Compare all three modes before choosing a deployment default.

Context mode comparison (--compare-context): runs the same live suite three times per model (Compact → Balanced → Full evidence) and writes a **Context mode comparison** table with overall score, pass rate, prompt/completion/total tokens, and mean latency. Use this to see whether a smaller context budget saves tokens/time without hurting investigation scores on your hardware.

Context-size (Findings + tools + history). Do not judge a local model only by tokens/sec — check tool use and grounding as context grows:

Context	Purpose
8K	Minimum investigation workload
16K	Typical investigation
32K	Large Findings / multi-tool investigation
64K	Stress test, if supported

Practical comparison on a developer workstation:


Gemini 3.7 Flash / Gemini 3.5 Flash-Lite
vs
Qwen3.5 9B (shipped default)
Qwen3.8 27B

Does extra local capacity improve investigation quality enough to justify memory and latency?

Reproducibility and architecture

Each run should save timestamp, app version, dataset version, model ids, endpoint config, cases, prompts, tool calls/results, final responses, scores, and timing. A run ID (AI_Benchmark_#2026-08-13-001) keeps results comparable when model behavior drifts without a viewer code change.

--fail-under N can fail CI when a model drops below a threshold (live often wants 0 so HTTP errors still write the report).



Investigation Case

BTFViewer uses one **Investigation Case** model (btf-investigation-case): question, scope (trace / C1-Cn / tasks / cores), hypotheses with status (**supported** / **pos-**

sible / need evidence / rejected), evidence graph, coverage, falsification checks, conclusion, and validation. See [Implementation notes](#).


After each final assistant reply a host-side **validator** extracts `jump:TIME` and `Task[id]` claims and flags invented names or timestamps outside the cursor window. **Test connection** appends a **model capability** card (live chat / structured output / tool calling overlay on a 3B vs 7B+ heuristic). Headless eval:

```
make -C BTFViewer ai-test  
# or: python builds/btf_viewer.py ai-test --dataset tests/ai --fail-under 70
```

Host investigation modes (quick / diagnose / compare / optimize / report) still map onto existing templates internally; they are not a permanent chip row. Reach the same workflows via the dynamic shortcuts or **More templates....** Always confirm on the timeline.

Investigation planner

Host-side planner. **Cheapest evidence first.** User-facing loop: [README](#) → [Investigation planner](#).




Tool / helper	Host behaviour
plan_investigation	Rank hypotheses and a cheap tool sequence from findings + question
suggest_scope	Task, related tasks, evidence times (or current cursors)
detect_contradictions	SUPPORTED / CONTRADICTED / INSUFFICIENT (e.g. execution $\gg$ blocking vs mutex hypothesis)
assess_evidence_sufficiency	Coverage heuristic $\rightarrow$ stop / continue / revise
score_hypotheses	Evidence-weighted scores (not a GUI tool)
cluster_findings	Group by shared task or pattern
generate_fingerprint	HIGH / MEDIUM / LOW scheduling, sync, timing bands
find_similar_investigations	Jaccard-style match vs recorded experiment outcomes

Tool / helper	Host behaviour
regression_localize	A vs B deltas → task / region / likely mechanism
build_causal_chain	Edges tagged causal / correlated / temporal; disclaimer required
generate_experiment_plan	Ranked pin / contention / priority experiments
record_experiment_outcome	Persist an outcome in the viewer's investigation storage
score_investigation_metrics	Host-computed on close (no longer a model-facing tool): Phase 3 extras evidence_efficiency, investigation_cost, false_confidence, falsification_quality, scope_accuracy, stop_efficiency (also spread into score_benchmark_case, with adversarial rates)

Do **not** add chat templates after auto_investigate.

Causal and temporal engines

Host-side heuristics over Analysis Findings — not an RTOS scheduler replay. User-facing loop stays on **Investigation planner**. Diagnose / Investigate / Auto investigate walk explanation tools before experiments: graph → temporal → rank_root_causes → challenge_conclusion, then what_if.



Tool / helper	Host behaviour
<code>analyze_temporal_causality</code>	Happens-before chain from finding times (jump:TIME)
<code>build_task_dependency_graph</code>	BTF sync/preempt/migrate/PI graph (finding-wording fallback); optional task neighborhood
<code>decompose_response_time</code>	Relative delay shares (mutex, preemption, migration, execution, scheduler)
<code>rank_root_causes</code>	Rank hypotheses or finding buckets
<code>verify_claim</code>	SUPPORTED / PARTIAL / UNSUPPORTED vs findings and cursors
<code>challenge_conclusion</code>	Alternatives and missing evidence
<code>investigation_memory</code>	Store or recall cases from the viewer's investigation storage
<code>cluster_incidents</code>	Group findings by time proximity

Tool / helper	Host behaviour
close_investigation	Record conclusion and close the case
analyze_distribution	p50 / p90 / p95 / p99 / p99.9, stddev, CV, 3-sigma outlier rate; BTF execution/blocking/PI/tick harvest
analyze_periodicity	Period/jitter from tick, STI, ISR, timer, or task-release times; kind = drift vs jitter vs WCET vs scheduler
summarize_investigation_context	Compact findings, hypotheses, and tools run

Engine limits

Engine	What it is	What it is not
analyze_temporal_causality	Happens-before from finding jump:TIME	Kernel event replay
build_task_dependency_graph	BTF sync / preempt / migrate / PI edges; 2-hop task neighborhood	Full ISR / object graph
decompose_response_time	Relative shares from finding magnitudes	Cycle-accurate milliseconds
rank_root_causes	Hypothesis or finding-bucket rank	A probability
investigation_memory	Local store / recall notepad	Team knowledge base
cluster_incidents	Time-proximity groups	Shared-mutex / causal clustering
close_investigation	Case status closed plus conclusion	Full firmware A/B lifecycle

Engine	What it is	What it is not
analyze_distribution	BTF execution / blocking / PI / tick samples (cap 8000)	A response-time series the parser does not have
analyze_periodicity	Inter-arrival jitter and kind	A kernel period timer
simulate_schedule	LEVEL 1 helper inside what_if	A GUI tool or RTOS kernel

Out of scope (do **not** add chat templates for these): trace-to-code (ELF / DWARF), real scheduler or hardware-aware simulation, model routing, automatic benchmark-case generation, natural-language → metric compiler, anomaly discovery without Analysis Findings, shared team investigation database.


Do **not** add chat templates after auto_investigate. Next gains come from deeper engines, not more buttons.

Implementation notes

Technical notes for the shared viewer AI implementation. For user-facing Case and Evidence behavior, see [README → Investigation Case](#). For live-suite XML, see [Benchmark and evaluation suite](#). Recorded scores are in [AI_BENCHMARK.md](#).

Analysis vs AI tools

Facts come from BTF Statistics pages first. AI ranks, explains, and navigates those facts. **Do not add AI tools** for work those pages already do (no detect_timeline_anomalies, extra jitter tools, or histogram tools). **Do not** invent kernel response time, inspect ELF/source, or simulate the scheduler.



The shipped loop stays **Triage → Investigate → Verify → Correlate → Critical Path → Dependency Graph → Temporal Causality → Rank → Challenge → What-if → Report**. Improve access to Statistics evidence (that is what `open_statistics_section` is for); prefer trimming the tool list over growing it — `score_investigation` and `investigation_replay` were folded into host-side close-out. User-facing page map: [README](#) → [BTF analysis pages](#).

Shared Case / Evidence engines


The viewer uses one Case, Evidence, planner, causal, tool, and Mermaid implementation. After AI UI changes, rebuild the distributed viewer artifacts with the project build targets.

UI lockstep: Start Investigation is separate. The template row shows up to five ranked shortcuts plus **More templates...** (flex-wrap, at most two lines). Suggested primary is outlined; Compare/SMP prerequisites show inline. Chip min-height 28px, disabled chips/menu items #8a96a8. Findings **Investigate...** uses the same outline style as the other Analysis footer buttons (not accent/primary). **More templates** use the same groups in a 2-column overlay. Findings **Save recipe...** and **Story...**

stay on that dialog. Trace Compare opens from toolbar **Compare**, not the Statistics footer. Composer **Context** row is collapsed by default; Apply cards label Navigation/Scope/Filter/Annotation/Export/Calculation; **Undo** restores Scope and Filters.

The ai-test CLI and offline benchmark use the same tests/ai fixtures (tracked .btf stubs + dataset.json). Live runs accept --context-mode and --compare-context (see [Context mode benchmarking](#)).

Validator



The host validator runs after the final reply. Prompting still forbids inventing numbers, task names, and jump:TIME; the validator is the guard, not the prompt.

Experiment close-out

validate_experiment compares expected vs actual signed percents (VALIDATED / PARTIALLY VALIDATED / DISPROVED), updates open hypotheses, and offers **Save to knowledge** (btfexp:save). Empty actual is filled from the last Trace Compare refresh (including **Limit to C1-Cn** when each tab has 2+ cursors) or compare_performance via experiment_percent_from_compare. Toolbar **Compare** → **Validate experiment...** closes the dialog and asks the model to call the tool with actuals omitted. Firmware-change capture remains a user step (new trace + toolbar **Compare**).

Capability, cost, privacy, knowledge

Feature	Host behaviour
Capability probe	Test connection lists models, chats with a JSON structured-output probe, then tool-calling (btf_ping then btf_pong). Live results overlay chat / structured output / tool calling / multi-tool chaining; long context and reasoning stay heuristic.
Cost	A dedicated usage bar shows Context: Compact · 4.6k tok · 3 tools · 12s (mode, tokens, tools, model time). Evidence uses the full format_cost_meter line. Clear resets replies, the meter, and current investigation issues.
Privacy	Chip ● Local / ● Cloud / ● Sensitive. Cloud send is blocked when sensitive; otherwise annotations are sanitized and optional task-name aliases apply (apply_cloud_privacy).
Knowledge	investigate matches user-saved entries (More → Save current finding...), then baseline, then the builtin catalog. Typical vs current rates show when both exist.
Interpret	Free-form Ask host-interprets (interpret_query), shows the scope card, then auto-runs (same as Run investigation). Templates / modes / prior assistant replies / short follow-ups skip the host interpret step. Scope toggles still allow re-run.
Tools used	Evidence Investigation details lists each tool with its host-side reason inline (no separate Why? action).

Diagrams

Replies may include Mermaid sequence diagrams for mutex, blocking, and priority events, or flowcharts for core migrations. **Compact** context mode emits diagrams only when the user asks. Pipe **Markdown tables** and sanitized HTML tables from Findings render as tables in the reply pane. In-chat Markdown / HTML tables follow the current theme. The Evidence panel also generates an Investigation tree when

investigate returns a root-cause chain. Diagrams follow the current light or dark theme; **Save As...** HTML exports use the light palette.

- Click a **task** node to lock-highlight that timeline row (Low[266] (Core 0) resolves to Low[266]).
 - Click a **core** node (Core_0, C0, C1) to switch to Core View and scroll to that core.
 - Mutex hex and other unresolved labels do nothing (the timeline stays undimmed).
 - Click empty figure area to open a larger zoom window (scroll to zoom 0.5-6×; **Esc** or **Close**). Trackpad pinch is treated as scroll.
 - The link row under the figure lists only jump:TIME and resolvable task/core names — not finding titles, hypotheses, or graph node ids (F, C0, ...).
 - **Save As...** HTML keeps inline SVG with clickable nodes (chat zoom wrappers are omitted).
-

Documentation navigation

Document	Question answered
README.md	How do I use BTFViewer?
WORKFLOWS.md	How do I diagnose a problem?
STATISTICS.md	What does this measurement mean?
AI.md	How does AI-assisted investigation work?
