


BTFViewer Beginner Workflow

This guide takes a first-time user through one complete investigation: learn the viewer controls, locate one task, read its Statistics, verify one incident on the timeline, and then use the AI Assistant to investigate and challenge the explanation.



The order matters. BTFViewer Statistics and the timeline provide measured evidence. **Analysis Findings** identify where to look. The AI Assistant organizes evidence and tests explanations. **What-if** and **Optimize** provide estimates only.

What you will learn

By the end of this workflow, you should be able to:

- distinguish Scope, Filter, Selection, and Highlight;
- use Task View, Core View, Load, Find, Cursors, Analysis, and Statistics;
- read Count, Avg, p95, p99, and Max without overinterpreting them;
- move from a Statistics value to the event that produced it;

- follow dependencies between timing, scheduling, and synchronization statistics;
- use AI after selecting evidence, not as the source of measurements; and
- validate a change with a new trace and equivalent measurements.

Running example: one task is sometimes late

This guide uses a generic task named `ControlTask`. Replace it with the task that matters in your trace.

The reported symptom is:

`ControlTask` usually behaves normally, but some activations finish later than expected.

The workflow does not assume the cause. Long response can come from longer own execution, preemption, an off-CPU wait, dispatch delay, synchronization, migration, or an incomplete trace. The purpose of the investigation is to distinguish these possibilities with evidence.

Essential terms

Term	Meaning in this workflow
Full Trace	The complete captured time span, with no cursor-defined analysis window
Scope	The time range used by Statistics, Analysis Findings, and AI: Full Trace or C1-Cn
Filter	A task, core, or migration subset inside the current Scope
Selection	The task or object kept active for inspection; it does not change calculations by itself
Highlight	Temporary visual emphasis; it does not change calculations
Slice	One continuous period when a task runs on one core
Off-CPU	Time when a task is not running; the trace may not show whether it was ready, blocked, suspended, or waiting for activation

Term	Meaning in this workflow
STI event	Optional software instrumentation for ready/resume, mutex, queue, interval, tag, lifecycle, or priority events
Finding	A deterministic or heuristic lead produced by Analysis; it is not a confirmed root cause
Hypothesis	A possible explanation that still requires supporting evidence and alternative checks

Selection and Highlight never silently become Filters. Always check the status bar and the Statistics **Filtered:** indicator before interpreting a value.

Evidence levels

Evidence level	Examples	How to use it
Direct	Recorded timestamp, task/core ID, slice boundary, STI tag	Strongest trace evidence
Derived	Execution duration, CPU share, migration count, inter-arrival	Deterministic calculation from recorded events
Estimated / heuristic	Response Time, mutex waiter-owner handoff, Critical Path, Task Health	Use to locate an episode, then confirm with stronger evidence
Configured comparison	Deadline or CPU-budget violation	Valid only when the configured threshold matches the real requirement
AI interpretation	Proposed cause, explanation, or recommendation	Verify against Statistics and the timeline

Do not describe an estimated value as a directly recorded kernel event.

Step 1 — Open and orient the trace

Actions

1. Select **Open** and load the trace.
2. Select **Fit Trace** (Ctrl+0 or F) so the complete capture is visible.
3. Start in **Task View**. Identify startup, steady-state, overload, idle, and shut-down phases.
4. Switch to **Core View** to see which task ran on each core.
5. Enable **Load** and resize the load chart if needed.
6. Hover over several timeline slices and read the task name, core, start/end time, and duration.
7. Confirm the status bar shows the expected active trace, **Scope: Full Trace**, Filters, and Zoom.

What to look for

- Does the trace contain the reported workload phase?
- Is ControlTask present and active?
- Are all expected cores visible?
- Are there long empty regions, activity bursts, or an unexpectedly short capture?
- Does the load graph show a distinct phase where the symptom may occur?

Continue when

You can identify the task of interest and the workload phase that should contain the problem. Do not choose a root cause yet.

Step 2 — Learn the essential tools

You do not need every toolbar function for a first investigation. Learn these controls first:

Tool	First-time use	What it changes
Task View / Core View	Follow one task or inspect per-core scheduling	Presentation only; Scope and Filters remain
Load	Show utilization over time	Adds the load chart; calculations do not change
Fit Trace	Return to the complete capture	Viewport only
Fit Cursors	Display the earliest-latest cursor span	Viewport only
Find (Ctrl+F)	Locate a task, migration, STI event, interval, lifecycle event, or object pointer	Moves between matches; Scope and Filters remain
Cursor	Mark an evidence time or define C1-Cn	Becomes Scope only when Limit to C1-Cn is enabled
Statistics	Calculate and display measurements	Uses the current Scope and active Filters
Analysis	Rank findings for the current Scope	Provides leads and navigation
Migration & Corridor Inspector	Inspect repeated core-to-core movement	Three columns: core path, heatmap, Topology. Click the heatmap for Path info. Analysis Scope defaults to Follow zoom (Fit = Full Trace, zoomed-in = Viewport). Lock Full Trace or Viewport, or choose Cursor C1-Cn

Tool	First-time use	What it changes
Compare	Compare Baseline and Candidate traces	Requires at least two open traces
AI Assistant	Explain, investigate, and verify supplied evidence	May request read-only or viewer-changing tools

Practice before analysis

1. Click the ControlTask label to make it the **Selection**.
2. Hover another task to see a temporary **Highlight**.
3. Open **Find**, search for ControlTask, and move through several matches with F3 and Shift+F3.
4. Place one cursor and move it by dragging.
5. Clear the cursor with the context menu or Shift+C.
6. Switch Task/Core View and confirm that Scope and Selection are preserved.

Continue when

You understand which actions only change the view and which actions change the analyzed data.

Step 3 — Check trace quality and the full-trace overview

Open **Statistics** in the right-side panel. Start with the default categories in this order:

OVERVIEW → TRIAGE → TIMING → SCHED → SYNC → DETAIL

For the first quality check, use only these sections:

Statistics section	Question to answer	Warning
Core Utilisation	Which cores were busy, idle, or unbalanced?	Average balance can hide a short overloaded phase

Statistics section	Question to answer	Warning
Trace Health (TICK)	Are TICK intervals, large gaps, and tickless behavior plausible?	Tickless idle can legitimately produce uneven intervals
Task Health	Which tasks deserve attention first?	The score is heuristic, not an AI probability
Core Time Breakdown	Does active, IDLE, TICK, and gap time account for the Scope?	Gap time is not automatically idle time
Top Tasks by CPU	Which tasks dominate on-CPU time?	CPU% uses wall-clock Scope duration and may total above 100% on multicore systems

Check whether the trace contains the STI data needed for the question. For example:

- dispatch latency needs a recorded ready/resume or create event;
- mutex, semaphore, and queue analysis needs corresponding STI events;
- explicit end-to-end operation time needs an Interval or other defined boundary;
- deadline results need a correct configured threshold.

If a required event is missing, record the limitation. Do not replace missing evidence with a precise explanation.

Continue when

The relevant workload phase is present, task/core data is plausible, and the trace has enough instrumentation for the intended analysis. Otherwise, fix the capture and record again.

Step 4 — Locate one task and one symptom

Keep the first investigation narrow. For the running example, the subject is `ControlTask`, and the symptom is occasional late completion.

Actions

1. Use **Find** or the Legend to locate `ControlTask`.
2. Select it so its timeline row remains easy to follow.
3. In **Statistics**, open **Response Time**, **Execution Time Per Slice**, **Blocking Time**, and **Period / Jitter**.
4. Do not apply a task Filter yet unless unrelated tasks make the data unreadable. Other tasks may explain preemption or synchronization delay.
5. Record the current Scope and any active Filters.

Translate the symptom into a measurement question

User report	First measurement question
"The task finishes late"	Is Response Time high, or is there an explicit Interval for the complete job?
"The task runs too long"	Is Execution Time per slice high?
"The task starts late after becoming ready"	Is Dispatch Latency high, and is a ready event available?
"The task wakes irregularly"	Are Period / Jitter and Inter-Arrival unstable?
"The task waits for a lock"	Do mutex STI events and Mutex Blocking support that statement?
"The task moves too much"	Are Migration Rate, Dwell, Ping, and core-pair paths unusual for this workload?

Use the user report to select the first statistic, not to declare the cause.

Continue when

You can state one measurable question, such as: "Which `ControlTask` samples have the longest response, and is the delay mainly own execution or off-CPU time?"

Step 5 — Read the Statistics correctly

Start with sample count

Always read **Count** or **Runs** before a percentile. A p99 calculated from ten samples is effectively the maximum and does not describe a stable one-percent tail.

Value	Beginner interpretation	Use
Count / Runs	Number of valid samples	Decide whether the result has enough evidence
Avg	Arithmetic mean	Describe the center, but not the tail
Median / p50	Middle ordered sample	Describe a typical sample with less outlier sensitivity
p95	95% of samples are at or below this value	Identify recurring slow-tail behavior
p99 / p99.9	99% / 99.9% are at or below this value	Inspect rarer tail delays when the sample count is sufficient
Max	Largest value observed in this capture and Scope	Jump to the worst observed event; not proven WCET
Jitter / Spread	Max – Min	Show the observed timing range
CV	Standard deviation divided by Avg	Compare relative variation across metrics with different scales

Read the shape, not only one number

For ControlTask:

1. Compare Avg, p95, p99, and Max.
2. Open the distribution or **Distribution Explorer**.
3. Use Scatter to see when spikes occur.
4. Use Histogram to see common value ranges.
5. Use CDF to see what percentage of samples are at or below a value.

Interpret common patterns cautiously:

Pattern	Possible meaning	Next check
Avg normal, Max much higher	One rare episode or small outlier group	Jump to Max and inspect its timeline window
p95 and p99 both high	Repeated tail delay	Check Recurring Patterns and several evidence times
Execution high with little off-CPU time	More own CPU work or long uninterrupted slices	Distribution, Intervals, Tags, preemption context
Response high but Execution normal	Delay is likely outside own on-CPU slices	Blocking, Dispatch, Preemption, Mutex, Migration
High CV	Behavior varies relative to its average	Scatter plot and workload-phase split

Continue when

You have selected one measured sample or recurring pattern that deserves timeline inspection.

Trace quality banner

When capture metadata reports overflow, truncation, or missing instrumentation, a **trace-quality banner** appears with **Review details**, **Continue with limitations**, and **Open capture guidance** (WORKFLOWS). Grouped details list affected Statistics and AI conclusions.

Guided first review

The guided demo (toolbar or Help) walks a sample trace end-to-end.

Step 6 — Use Analysis Findings for triage

Select **Analysis**. Findings are calculated for the current Scope and stay open while the timeline remains interactive.

For each relevant finding:

1. Read its severity, title, measured value, and **Evidence** line.
2. Treat severity as attention priority, not failure probability.
3. Select **Show on timeline** to center the timeline on the cited event without changing Scope or Filters.
4. Select **Investigate** to open the supporting Statistics section and apply a recommended cursor range when offered.
5. Confirm that the Statistics value, task name, and timestamp match the finding.
6. Mark **Done** when reviewed, **Dismiss...** with a short reason when not applicable, or **Add to case** to pin the finding on the AI Case.

If no relevant finding exists, continue from the measured Statistics sample. A missing finding does not mean the task has no problem; the symptom may not match a built-in heuristic.

Good first triage sections

Section	Use it for
Timeline Anomalies	Ranked unusual tails, bursts, gaps, migrations, and deadline events
Worst Events	Longest observed execution, blocking, inter-arrival, and response episodes
Recurring Patterns	Repeated anomaly types for the same task
Task Health	A broad heuristic screening score and links to component sections

Continue when

You have one task, one metric, and at least one evidence timestamp or interval.

Step 7 — Scope one incident with Cursors

Full-trace Statistics can mix startup, steady-state, overload, recovery, and shutdown. Use Cursors to turn one event into a testable incident.

Actions

1. Jump to Max, p95, p99, an anomaly, worst event, chart point, or finding evidence.
2. Place **C1** before the activity that may have triggered the delay.
3. Place **C2** after ControlTask completes or recovers. Additional cursors may mark intermediate evidence.
4. Select **Fit Cursors** (Ctrl+R) to display the earliest-latest cursor span.
5. Enable **Limit to C1-Cn** in Statistics, from the toolbar **C1-Cn** chip, or from the banner: **Use C1-Cn as analysis Scope** → **Enable Limit to C1-Cn**. The chip uses the Scope colour when Limit is on.
6. Confirm the status bar and Statistics header show **Scope: C1-Cn · duration**.
7. Recheck the **Filtered:** indicator. Clear unintended task, core, or migration Filters.
8. Reopen **Analysis** so findings use the same Scope.
9. Add a Bookmark or Annotation at the strongest evidence time.

Choosing a useful window

- Start early enough to include the possible trigger, previous task slice, ready event, mutex acquisition, or migration.
- End late enough to include completion, release, or recovery.
- Avoid a window so wide that unrelated phases dominate the measurements.
- Avoid a window so narrow that the cause lies outside it.

Navigation, Scope, and Filter are different. **Show Evidence**, Find, and Fit change where you look. **Limit to C1-Cn** changes which samples are calculated.


Continue when

The unusual value still appears inside the cursor Scope and you can see the surrounding task/core activity.

Step 8 — Follow statistic dependencies


Do not open every table. Start with the symptom and follow the smallest dependency path that can confirm or reject a cause.

Path A — task latency



Primary statistic	What it measures	Confirm with	Important limit
Response Time	Heuristic previous-slice-end to current-slice-end window	Execution, Blocking, Dispatch, Interval	Not an explicit release-to- completion pair
Execution Time	One continuous on-CPU slice	Distribution, Interval, Tags, Preemption	A task job may contain several slices
Blocking Time	Off-CPU gap between consecutive slices	Preemption, Mutex Blocking, Period	Off-CPU does not prove mutex blocking
Dispatch Latency	Recorded ready/create event to next switch-in	Timeline, Preemption, Core Time	Requires suitable STI ready evidence
Critical Path	Heuristic ready-to-completion window with overlapping evidence	Execution, Blocking, Preemption, Migration	Components can overlap; they are not a stacked duration split
Interval Analysis	Explicit instrumented start-to-stop duration	Timeline, Execution, Tags	Correctness depends on application instrumentation

For `ControlTask`, first decide whether the late response contains unusually long own execution or mostly off-CPU time. That decision determines the next path.

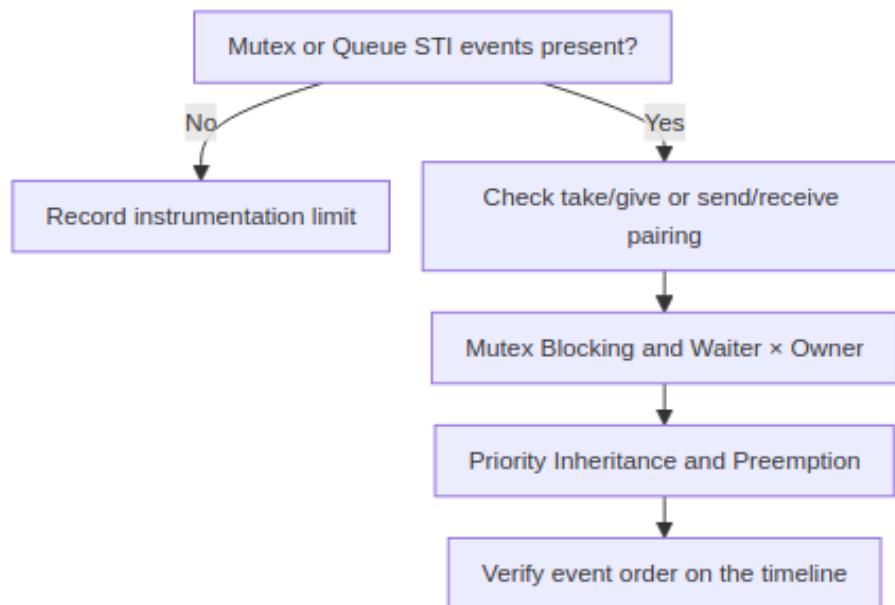
Path B — scheduling and multicore placement

Check load balance before calling migration a problem. An SMP scheduler may move an unpinned task to use an idle core.

Question	Statistics
Is imbalance persistent or phase-specific?	Core Utilisation, Core Utilization Over Time
Which task contributes to each core?	Task × Core
How often and how quickly does the task move?	Core Migrations: Count, Rate, Dwell, Ping
Which directed path dominates?	Core-Pair Migration Summary
Was the placement allowed?	Core Affinity
Did the move coincide with delay or scheduling cost?	Response, Execution, Switch Overhead, Preemption

Migration count alone does not measure cache misses, coprocessor-save cost, or scheduler overhead. Use processor-specific evidence when making those claims.

Path C — synchronization and waiting



Use **Mutex / Semaphore** or **Queue** to check event-pair quality before trusting derived waits. **Waiter x Owner** and **Mutex Blocking** use heuristic handoffs; they do not reconstruct the RTOS wait queue. Confirm the object pointer, owner, waiter, take/give order, and priority activity on the timeline.

Path D — periodicity and jitter

Start with	Check next	Question
Period / Jitter	Inter-Arrival Time	Are activations missed, extra, or bursty?
Inter-Arrival Time	Unified Jitter	Is variation in activation spacing recurring?
Unified Jitter	Execution, Blocking, Response, Dispatch	Which component contributes most variation?


Start with	Check next	Question
Core Utilization Over Time	Preemption, Mutex, Migration	Does a load burst coincide with the timing change?
Deadlines / CPU Budget	Execution, Period, Critical Path	Does the measured behavior violate a real configured requirement?

Continue when

The leading explanation is supported by a primary statistic and at least one dependent statistic, or the evidence shows that another path is more plausible.

Step 9 — Verify the event on the timeline

Statistics summarize samples. The timeline shows event order. A useful conclusion needs both.



Ask these questions:

1. Does the value reproduce inside the current C1-Cn Scope?
2. Does clicking the value open the expected task, core, and time?
3. What ran immediately before, during, and after the delay?
4. Is ControlTask running, ready, preempted, blocked, suspended, or waiting for its next activation? Does the trace contain enough events to tell?
5. Do mutex, queue, priority, migration, or ready events support the proposed mechanism?
6. Is the relationship causal, correlated, or only close in time?
7. What evidence would disprove the explanation?
8. Can a simpler explanation fit the same evidence?

Confidence labels

Label	Minimum evidence
Supported	Metric reproduces, exact timeline evidence matches, and reasonable alternatives were checked
Plausible	Some evidence matches, but an important event or relationship is missing
Inconclusive	The trace cannot distinguish the main alternatives
Unsupported	The scoped metric or timeline contradicts the explanation

Use **Confirmed** only when the application requirement and instrumentation define the relevant boundaries, and repeated equivalent measurements support the same conclusion.

Continue when

You can write one evidence-based statement without using the AI, for example: “The longest observed ControlTask response in this Scope contains a normal execution slice but a long off-CPU interval that overlaps specific preemption activity.”

Step 10 — Use the AI Assistant

AI is optional. Use it after selecting a finding, task, event, distribution, or C1-Cn range. It receives structured Findings and tool results, not the complete raw .btf event stream.

One-time setup


1. Open **Settings → AI**.
2. Select a provider preset, endpoint, model, and authentication method.
3. Use **Test connection**.
4. Start with **Balanced** Context. Use Compact for a smaller context window or Full evidence when the investigation needs more findings and history.

5. Review privacy settings before using a cloud endpoint. Local Ollama normally needs no API key.

Choose the entry point that matches the evidence

Evidence already selected	AI entry point	Context sent
No clear starting issue	Start Investigation or Triage findings	Current Scope, Filters, and available findings
One Analysis Finding	Investigate, Explain, Verify, or Auto investigate	Selected finding and its evidence
C1-Cn incident	Explain this region with AI or Explain region	Cursor range and scoped findings
One timeline segment	Ask AI about this event	Selected task, core, segment, and nearby evidence
Open distribution	Query with AI...	Selected metric, task, and displayed samples
Migration & Corridor Inspector	Investigate with AI	Analysis scope, selected path, ping-pong/dwell, handoff heuristic, load balance, Inspector filters
Two comparable traces	Trace Compare Query with AI...	Selected comparison tables

Follow the six AI stages



For the running example:

1. Ask **Investigate** to explain the scoped ControlTask incident.
2. Open every cited Statistics section and jump:TIME / range:L0/HI link.
3. Check that every cited time lies inside C1-Cn.
4. Use **Verify finding** to request supporting evidence, contradictory evidence, alternatives, and missing information.
5. Reject any statement that cannot be reproduced in Statistics or assumes an unrecorded event.
6. In **Evidence & Validation**, click **[Run]** on ► **Next check** to continue in the same Investigation Case, Context, and Scope. Extra host follow-ups are under **More next steps....**
7. If the reply includes a dedicated nextstep:{action} line (nextstep: in English; the action in the reply language), that line also has **[Run]** and sends that sentence. A heading such as **Next check:** without that tag is not a button.
8. Confirm the follow-up on the timeline and named Statistics pages before treat-

ing it as done.

Understand AI tool actions

Tool behavior	What happens
Evidence query	Runs immediately and returns measured or derived data
Investigation-state or export tool	Runs immediately; may update hypotheses, memory, experiment records, or save a report
Viewer-changing tool	Waits for Apply or Skip unless Auto-apply GUI actions is enabled

Viewer-changing tools can place Cursors, zoom, highlight a task, change View Mode, open the Corridor Inspector, add a note or bookmark, clear marks, or reset the view. Inspect the tool card before selecting **Apply**. Use **Undo last actions** if an applied batch is not useful.

A useful AI answer should include

- the active Scope and focus task;
- measured observations separated from interpretation;
- supporting and contradictory evidence;
- alternative explanations;
- a verdict such as Supported, Rejected, or Inconclusive;
- missing evidence and a tagged `nextstep:{action}` follow-up (nextstep: stays English; the action uses the reply language); and
- a clear estimate disclaimer for What-if or Optimize.

Continue when

The AI explanation matches the Statistics and timeline evidence, or it has identified a specific missing measurement that requires another capture.

Step 11 — Define one change and measure it again

Do not move from a Finding directly to a fix. Make a change only after the evidence supports a mechanism.

Actions

1. Define one change, such as task affinity, priority, mutex scope, workload distribution, or instrumentation.
2. Write the expected metric effect before applying the change.
3. Optionally use **What-if** or **Optimize** to rank experiments. Treat the output as a heuristic estimate, not an RTOS scheduler simulation.
4. Apply the change to the real system.
5. Capture the same workload again with equivalent instrumentation.
6. Open the original trace as **Baseline** and the new trace as **Candidate**.
7. Select equivalent workload phases and cursor ranges.
8. Repeat the same Statistics used in the original investigation.
9. Use **Compare** to review normalized totals, tail values, notable changes, and side effects.
10. Return to both timelines and inspect the samples behind the difference.

Example acceptance statements

- Response p99 decreases without increasing deadline misses.
- Blocking tail decreases without increasing migration rate or core imbalance.
- Migration rate decreases without overloading the selected core.
- Load balance improves while execution and response tails remain within the requirement.

If the result does not match the prediction, revise the hypothesis. Do not rewrite the explanation to fit the outcome.

Continue when

Baseline and Candidate represent equivalent conditions, the target metric has been remeasured, and important side effects have been checked.

Step 12 — Record and share the investigation

Record enough information for another engineer to reproduce the result:

Record	Include
Capture context	Trace names, firmware/configuration, workload, core count, instrumentation
Analysis context	Scope, Filters, selected task, and relevant cursor times
Symptom	User-visible problem and affected task/core
Measurements	Count, Avg, p95, p99, Max, units, and exact evidence times
Dependencies	Supporting timing, scheduling, synchronization, or instrumented statistics
Conclusion	Supported explanation, confidence, alternatives, and missing evidence
Experiment	Change, expected result, Candidate measurement, and side effects

Useful outputs include:

- Bookmarks and Annotations for important timestamps;
- an annotated Snapshot or Timeline SVG;
- **Save cursor range as BTF** for the selected incident;
- Statistics **Export HTML** (self-contained report with searchable tables);
- Trace Compare **Export HTML**; and
- AI diagnostic report or Investigation Case when AI was used.

Keep the source trace. An exported report summarizes evidence but cannot preserve every interactive timeline action.

Complete worked example

The following example shows the entire path without inventing numeric results.

Stage	Action for ControlTask	Decision
1	Open trace, Fit Trace, enable Load, identify steady-state phase	The reported workload is present
2	Select ControlTask, practice Find and Task/Core View	Task and surrounding cores are visible
3	Check Core Utilisation, Trace Health, and instrumentation	Trace is usable for timing analysis
4	Open Response, Execution, Blocking, and Period/Jitter	The question becomes “own execution or off-CPU delay?”
5	Compare Count, Avg, p95, p99, Max, and distribution	One tail sample or recurring group is selected
6	Open Analysis; use Show Evidence or Investigate	Finding and Statistics point to the same episode
7	Place C1 before the trigger and C2 after completion; enable Limit	Statistics now describe one incident
8	Follow Response → Execution/Blocking → Preemption/Mutex/Migration	The smallest supporting dependency path is collected
9	Verify exact task/core/event order on the timeline	Explanation is Supported, Plausible, Inconclusive, or Unsupported
10	Ask AI to Investigate, then Verify finding; continue with Evidence [Run] or a tagged nextstep: {...} [Run]	AI explanation is checked against the same evidence
11	Define one expected metric change, capture Candidate, Compare	The change is measured rather than assumed
12	Save Scope, values, evidence times, conclusion, and report	Another engineer can reproduce the investigation

Beginner completion checklist

- ☐ I confirmed the active trace, Scope, Filters, and View Mode.
- ☐ I checked trace quality and required STI instrumentation.
- ☐ I selected one task and one measurable question.
- ☐ I read Count before interpreting p95 or p99.
- ☐ I compared Avg, p95, p99, Max, and the distribution.
- ☐ I treated Analysis Findings as leads, not conclusions.
- ☐ I placed at least two Cursors and enabled **Limit to C1-Cn**.
- ☐ I followed the relevant statistic dependencies.
- ☐ I opened the exact sample on the timeline.
- ☐ I checked contradictory evidence and alternative explanations.
- ☐ I used AI only after selecting evidence.
- ☐ I continued the AI investigation with Evidence **[Run]** or a conversation nextstep: {...} **[Run]**.
- ☐ I verified every AI measurement and timestamp.
- ☐ I treated What-if and Optimize as estimates.
- ☐ I captured an equivalent Candidate and repeated the same measurements.
- ☐ I recorded enough context for another engineer to reproduce the result.

Common beginner mistakes

Mistake	Better practice
Starting with AI on the entire trace	Select a Finding, task, event, distribution, or C1-Cn range first
Confusing Highlight with Filter	Check the status bar and Statistics Filtered: indicator
Treating viewport zoom as Scope	Enable Limit to C1-Cn and confirm the Scope label
Reading p99 from very few samples	Read Count first and treat a small-sample percentile cautiously
Calling Max a guaranteed WCET	Say “maximum observed in this capture and Scope”

Mistake	Better practice
Calling every off-CPU gap mutex blocking	Require synchronization evidence or keep the term Off-CPU / Blocking Time
Calling every migration harmful	Check load balance, affinity, rate, dwell, and correlated delay
Assuming two nearby events prove causation	Check event order, alternatives, and contradictory evidence
Comparing different workload phases	Match workload, instrumentation, Scope, and Filters
Treating an AI explanation as a measurement	Reproduce it in Statistics and on the timeline
Expecting Next check: in the reply to be a button	Use Evidence [Run] , or a dedicated nextstep:{action} line
Treating What-if as a verified improvement	Apply the change, capture again, and Compare

When to stop and recapture

Stop the current investigation and capture again when:

- the reported phase or incident is missing;
- the required task, core, TICK, ready, mutex, queue, interval, or tag events were not recorded;
- Scope boundaries cut off the event needed to match a pair;
- Baseline and Candidate do not represent equivalent workloads;
- the next useful check requires new instrumentation; or
- the evidence cannot distinguish the remaining hypotheses.

An inconclusive result with a clear instrumentation requirement is more useful than a confident explanation without evidence.

Documentation navigation

- [README.md](#) — installation, supported files, toolbar, timeline controls, export,

and shortcuts

- [STATISTICS.md](#) — definitions, calculations, dependencies, interpretation, and limitations for every Statistics section
- [AI.md](#) — AI setup, actions, tools, privacy, evidence validation, and advanced reference
- [WORKFLOWS_zh-TW.md](#) — Traditional Chinese version